








云原生社区  
Cloud Native Community

云原生社区 MEETUP

# 基于VK实现多集群管理

王琼 (YY直播 SRE)





## 集群管理

集群接管  
节点上下线  
离线资源池管理  
资源池管理  
驱逐



## 持续集成

代码构建  
镜像构建




## 服务发布

资源大小选择  
资源池选择  
调度策略选择  
固定ip支持  
回滚  
灰度  
配置管理



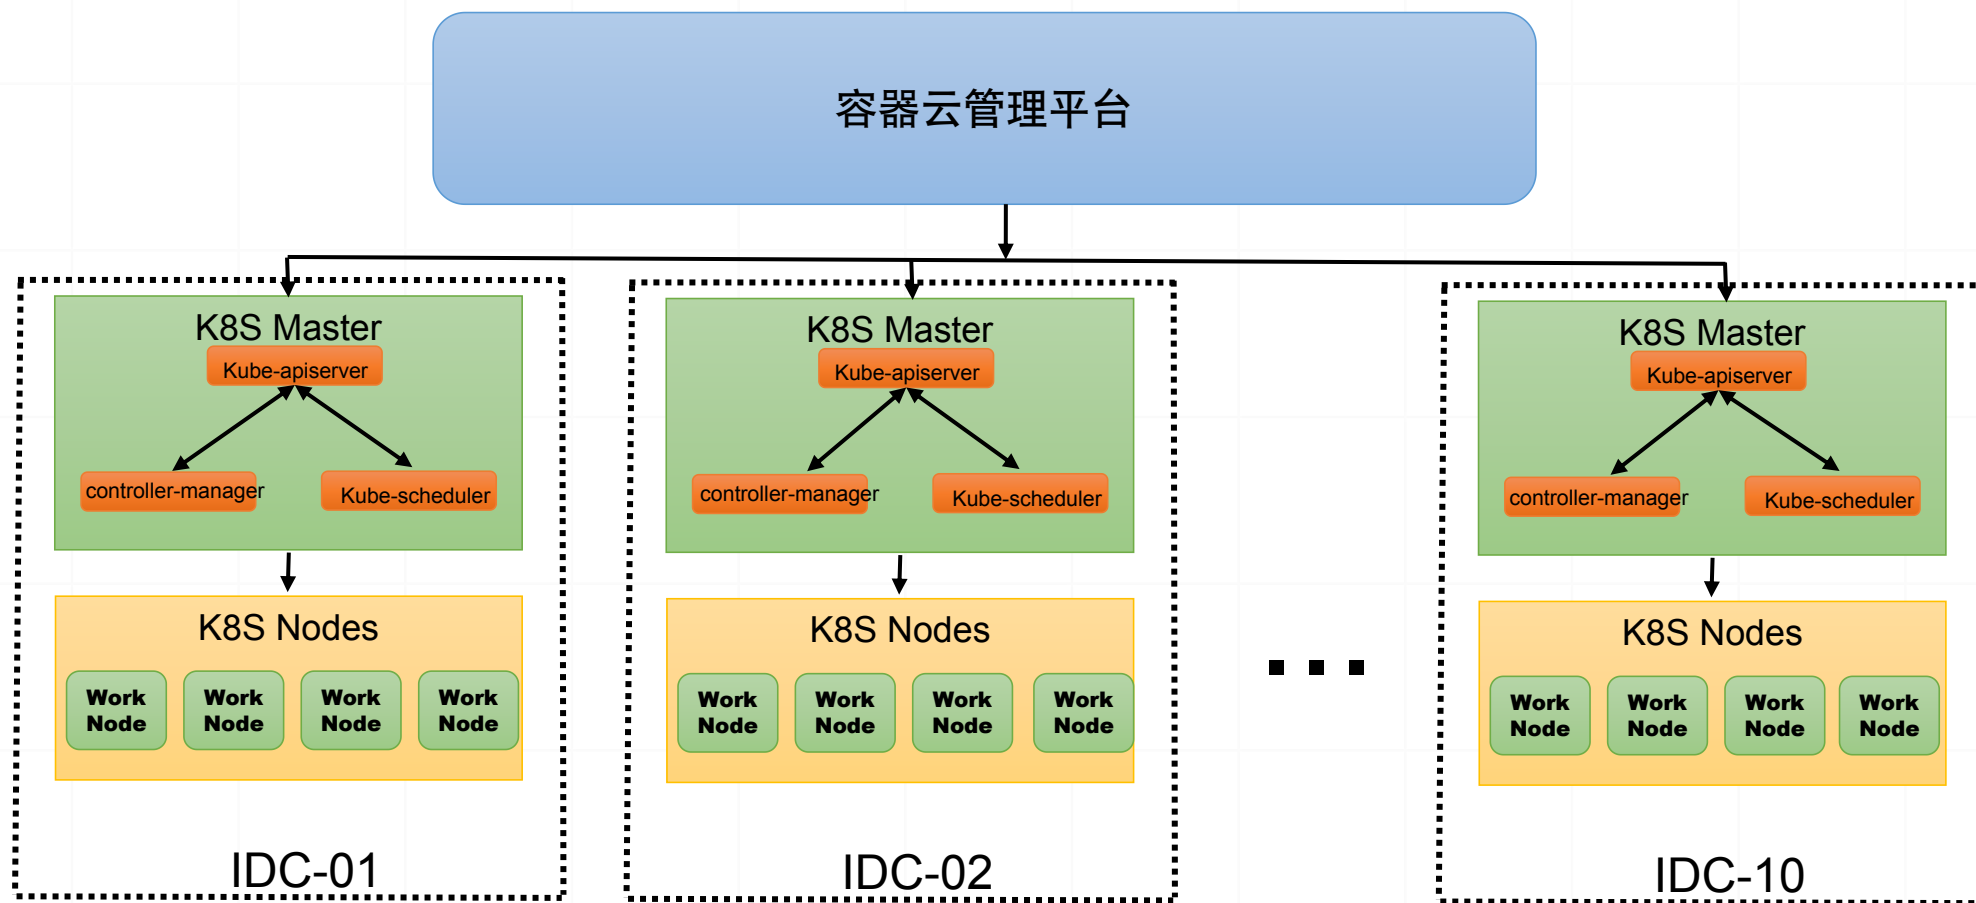
## API提供

提供统一的部署接口  
提供查询功能



## 报表查询

容器化接入率  
集群资源使用详情  
集群监控





## 单集群故障

单集群越来越大，  
降低单集群故障对  
业务的影响

## 集群平滑升级

通过新增集群实  
现集群的平滑升  
级

## 第三方集群接入

业务无需代码改  
动即可发布服务  
至第三方集群

## 边缘机房区域自治

防止前端机房计算节  
点与控制平面失联

## 公有云资源接入

接入公有云弹性资源  
实现快速弹性功能



## 解决的问题：

依靠云资源可秒级弹性扩容

无需预留离线资源

支持优先部署到自建机房，资源不足时部署到弹性机房

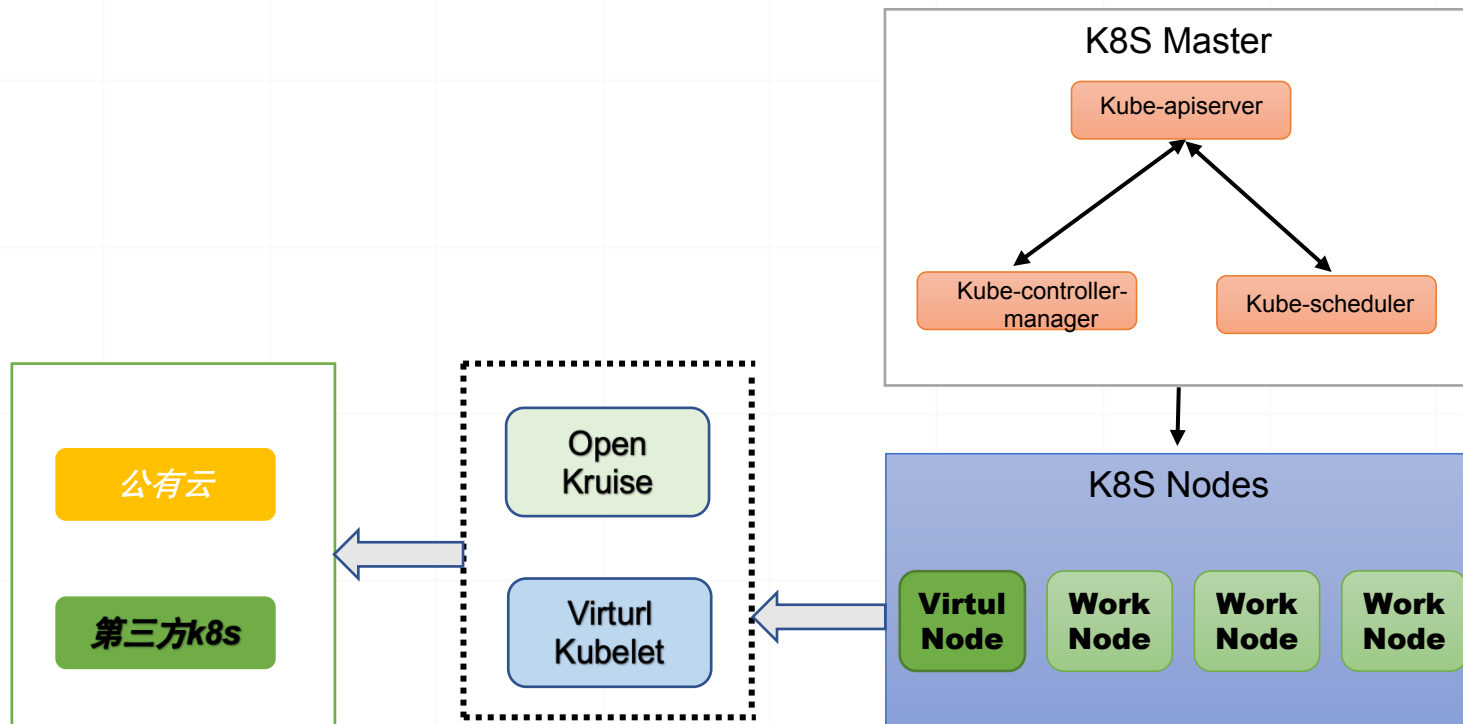
快速接入第三方集群



## 成本优化：

云厂商即用即计费，优化计算资源成本，按需计费，不使用不计费


抢占式实例成本优化明显





## 公有云费用对比

■ 包月 ■ 按量计费 ■ 按需计费



01

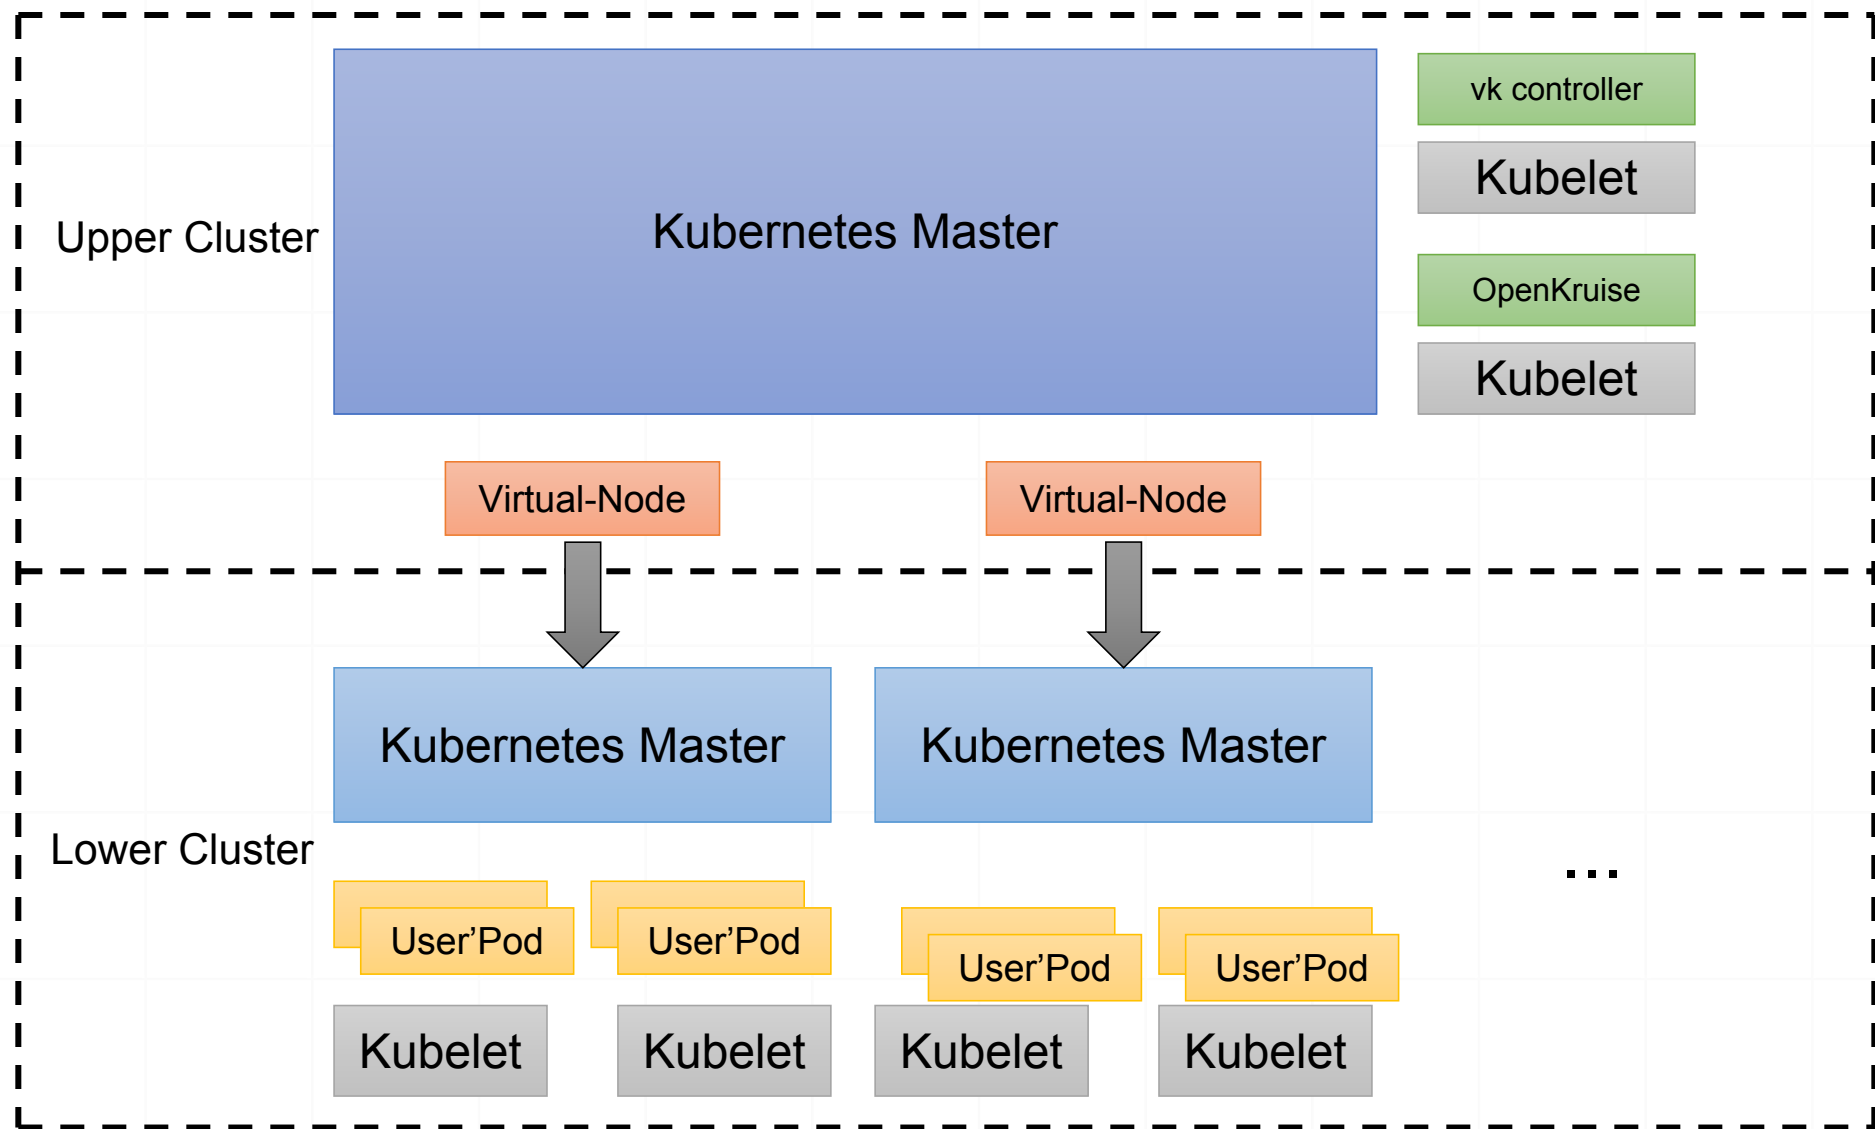
节点弹性伸缩在定时扩缩容节点满足目前的需求，扩容需要5分钟左右，在非定时伸缩情况下需要15分钟完成扩容

02

弹性伸缩可实现秒级扩容，与正常pod实例无差别

03

弹性伸缩比起节点弹性伸缩成本上表现更优，可节省更多的计算机缘成本





## Virturl Kubelet

- 虚拟 WorkerNode

Liqo Provider

Tensile Kube Provider



## OpenKruise

- WorkloadSpread

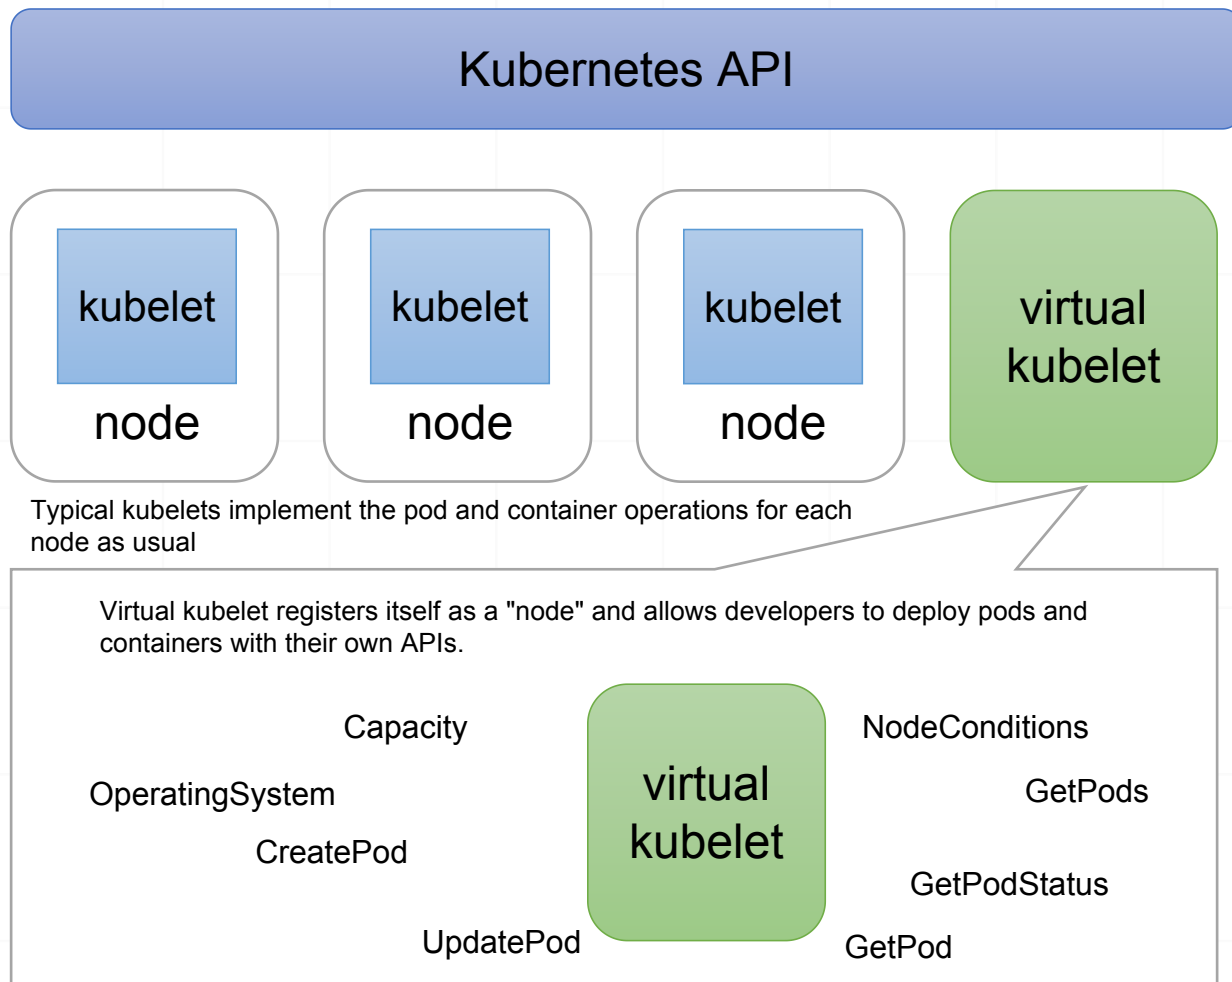
能够将 Workload 的 Pod 按一定规则分布到不同类型的 Node 节点上，  
赋予单一 Workload 多区域部署和弹性部署的能力

- SidecarSet

自动注入 Container 至 Pod

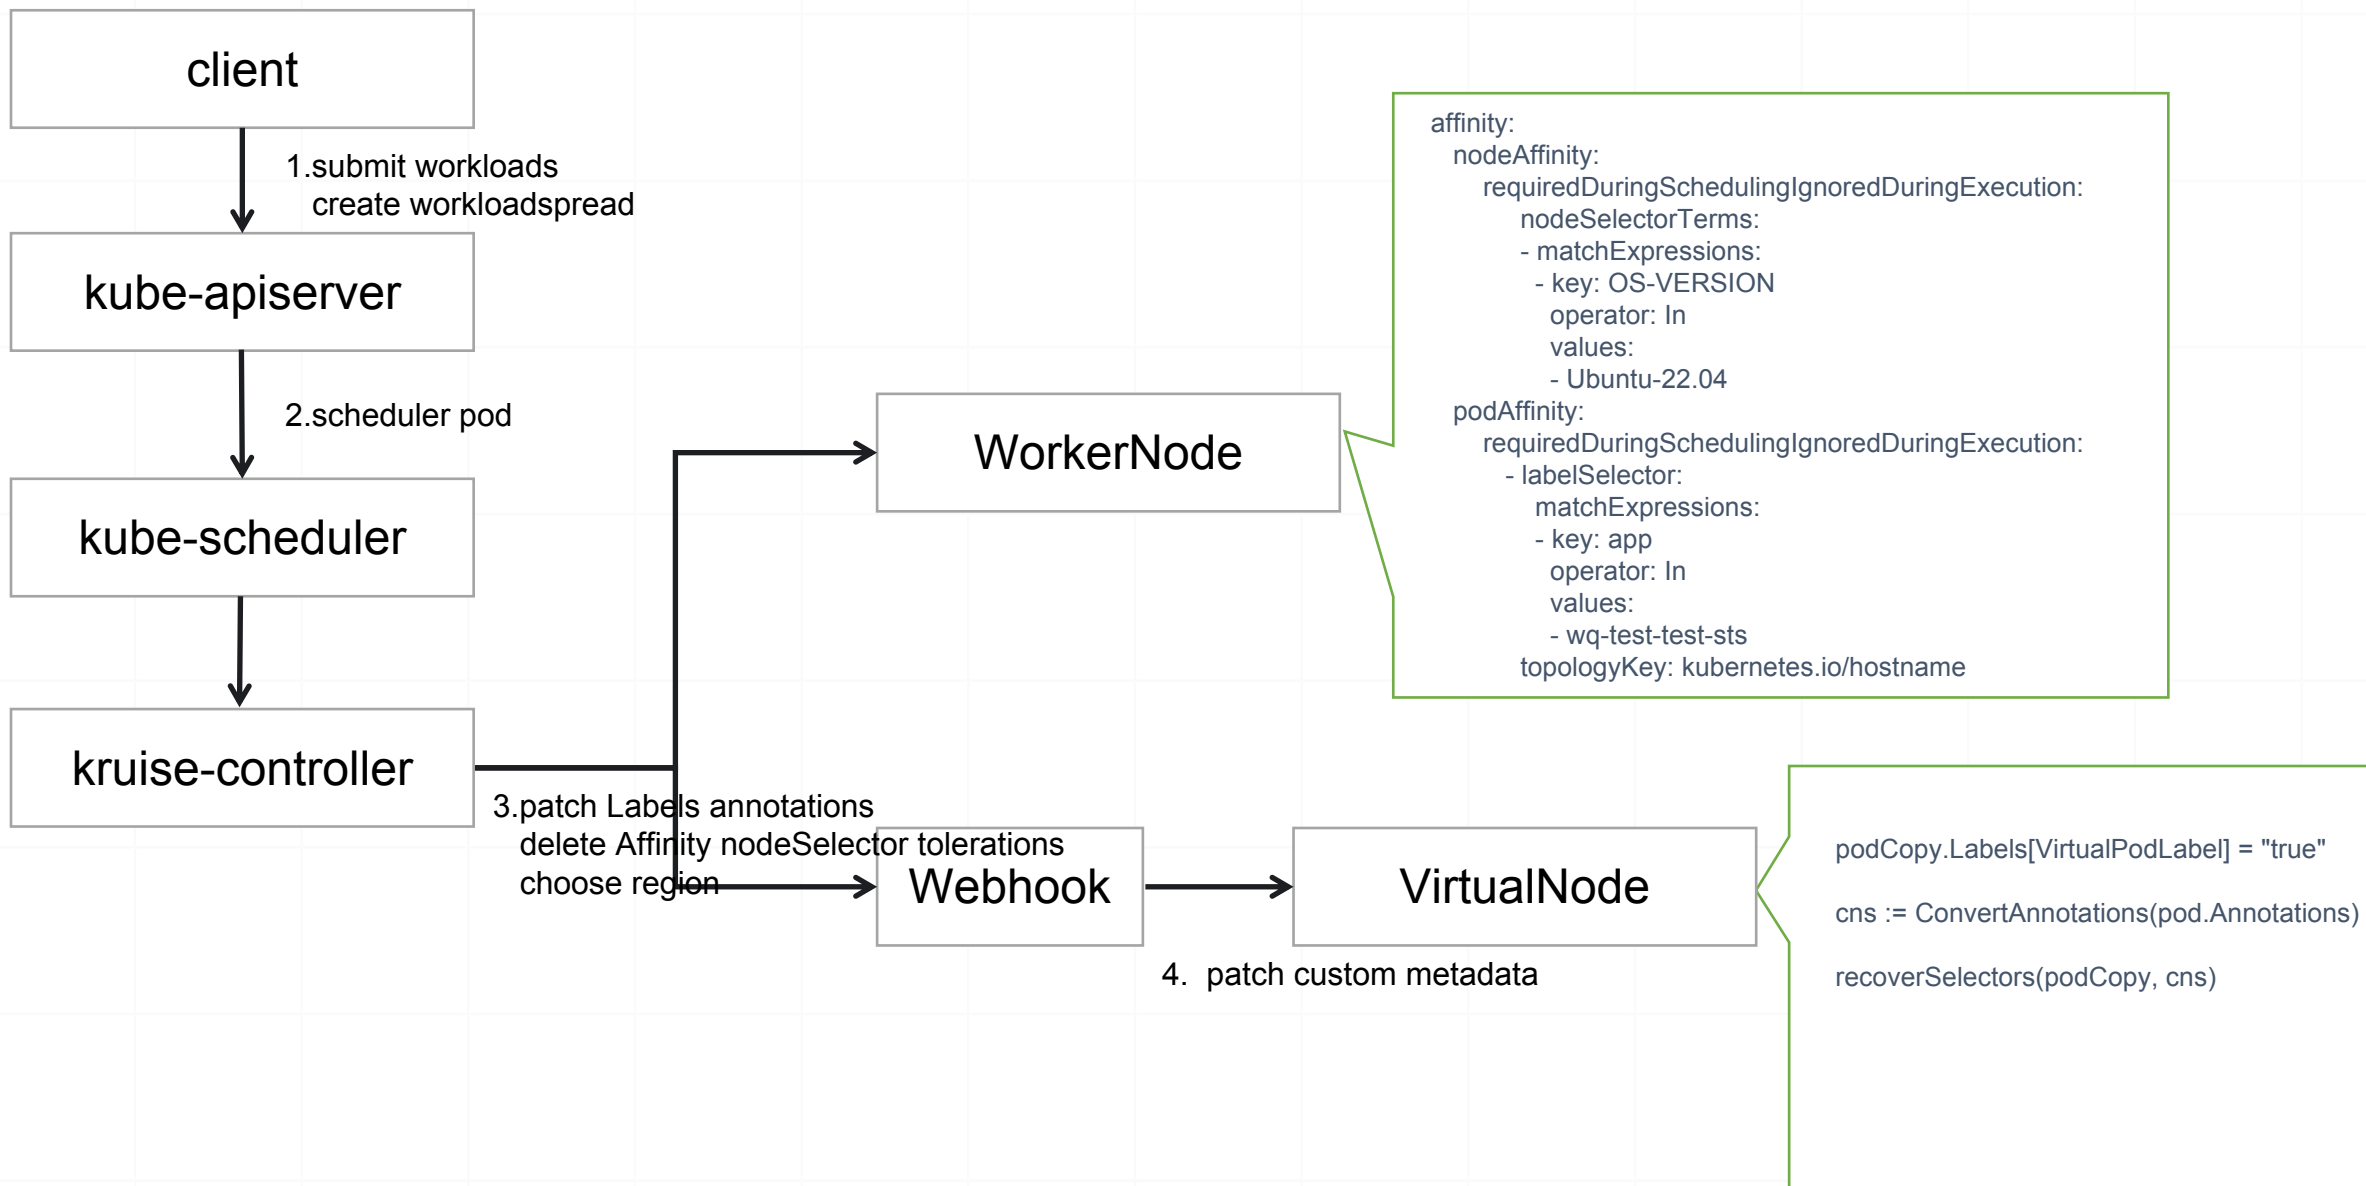
- Webhook





## 功能：

- create, delete and update pods,cm,secrets
- container logs, exec, and metrics
- get pod, pods and pod status
- capacity
- node addresses, node capacity, node daemon endpoints
- operating system
- bring your own virtual network





## WorkloadSpread

能够将 Workload 的 Pod 按一定规则分布到不同类型的 Node 节点上，赋予单一 Workload 多区域部署和弹性部署的能力

优先部署到自建机房，资源不足时部署到 VK

优先部署固定数量个 Pod 到自建机房，其余到 VK

```
apiVersion: apps.kruise.io/v1alpha1
kind: WorkloadSpread
metadata:
  name: workloadspreload-demo
spec:
  scheduleStrategy:
    adaptive:
      rescheduleCriticalSeconds: 30
    type: Adaptive
  subsets:
  - name: common
    maxReplicas: 50%
    patch:
      metadata:
        annotations:
          subset: common
      requiredNodeSelectorTerm:
        matchExpressions:
        - key: node-role.kubernetes.io/WorkerNode
          operator: In
          values:
          - "true"
      tolerations:
      - effect: NoSchedule
        key: unscheduer
        operator: Equal
        value: "true"
      - name: local-virtual-kubelet
        patch:
          metadata:
            annotations:
              clusterSelector: '{"tolerations":[{"key":"node.kubernetes.io/not-ready","operator":"Exists","effect":"NoExecute"}, {"key":"node.kubernetes.io/unreachable","operator":"Exists","effect":"NoExecute"}, {"key":"test","operator":"Exists","effect":"NoExecute"}, {"key":"test1","operator":"Exists","effect":"NoExecute"}, {"key":"test2","operator":"Exists","effect":"NoExecute"}, {"key":"test3","operator":"Exists","effect":"NoExecute"}]}'
            subset: local-virtual-kubelet
          spec:
            affinity:
              nodeAffinity: null
              podAntiAffinity: null
              nodeSelector: null
            requiredNodeSelectorTerm:
              matchExpressions:
              - key: node-role.kubernetes.io/local-virtual-kubelet
                operator: In
                values:
                - "true"
```



```
func recoverSelectors(pod *corev1.Pod, cns *ClustersNodeSelection) {
    if cns != nil {
        pod.Spec.NodeSelector = cns.NodeSelector
        pod.Spec.Tolerations = cns.Tolerations
        if pod.Spec.Affinity == nil {
            pod.Spec.Affinity = cns.Affinity
        } else {
            if cns.Affinity != nil && cns.Affinity.NodeAffinity != nil {
                if pod.Spec.Affinity.NodeAffinity != nil {
                    pod.Spec.Affinity.NodeAffinity.RequiredDuringSchedulingIgnoredDuringExecution =
cns.Affinity.NodeAffinity.RequiredDuringSchedulingIgnoredDuringExecution
                } else {
                    pod.Spec.Affinity.NodeAffinity = cns.Affinity.NodeAffinity
                }
            } else {
                pod.Spec.Affinity.NodeAffinity = nil
            }
            if cns.Affinity != nil {
                if cns.Affinity.PodAffinity != nil {
                    pod.Spec.Affinity.PodAffinity = cns.Affinity.PodAffinity
                }
                if cns.Affinity.PodAntiAffinity != nil {
                    pod.Spec.Affinity.PodAntiAffinity = cns.Affinity.PodAntiAffinity
                }
            }
        }
    } else {
        pod.Spec.NodeSelector = nil
        pod.Spec.Tolerations = nil
        if pod.Spec.Affinity != nil && pod.Spec.Affinity.NodeAffinity != nil {
            pod.Spec.Affinity.NodeAffinity.RequiredDuringSchedulingIgnoredDuringExecution = nil
        }
    }
    if pod.Spec.Affinity != nil {
        if pod.Spec.Affinity.NodeAffinity != nil {
            if pod.Spec.Affinity.NodeAffinity.RequiredDuringSchedulingIgnoredDuringExecution == nil &&
                pod.Spec.Affinity.NodeAffinity.PreferredDuringSchedulingIgnoredDuringExecution == nil {
                    pod.Spec.Affinity.NodeAffinity = nil
            }
        }
        if pod.Spec.Affinity.NodeAffinity == nil && pod.Spec.Affinity.PodAffinity == nil &&
            pod.Spec.Affinity.PodAntiAffinity == nil {
            pod.Spec.Affinity = nil
        }
    }
}
```




## SidecarSet

自动注入sidecar,可以将日志采集、监控上报等agent,通过sidecar的形式自动注入到pod

```
# sidecarset.yaml
apiVersion: apps.kruise.io/v1alpha1
kind: SidecarSet
metadata:
  name: test-sidecarset
spec:
  selector:
    matchLabels:
      vk: true
  updateStrategy:
    type: RollingUpdate
    maxUnavailable: 1
  containers:
    - name: sidecar1
      image: centos:6.7
      command: ["sleep", "999d"] # do nothing at all
      volumeMounts:
        - name: log-volume
          mountPath: /var/log
  volumes: # this field will be merged into pod.spec.volumes
    - name: log-volume
      emptyDir: {}
```




Service



PVC



CRD



感谢观看



云原生社区  
Cloud Native Community

