



OBSERVABILITY
SUMMIT 2023

可观测性峰会 第1届

基于 Prometheus 的 SLO 告警实战


宋佳洋@OPPO



大纲

1. SLO 告警基础知识
2. 基于 Sloth 项目构建
3. 基于 Pyrra 项目构建
4. 多租户 SLO 服务构建

关于我



先后就职于七牛云、京东云等公司，目前在 OPPO 从事云计算相关工作。



爱好开源，目前主要关注 Go 和 云原生可观测领域、是开源项目 Prometheus、Cortex、Thanos 的代码贡献值。



songjiayang



微信公众号：Grafana 爱好者

为什么基于 SLO 告警

~~100%~~



- 梳理内容
- 优先级告警
- 利益方认可
- 持续迭代



没有 SLO, 就没有 SRE

SLO 相关概念

SLI

- 状态码 ≥ 500
- 请求延迟 $> 200\text{ms}$
- 进程运行非 0 状态码退出

错误预算

- 时间周期：30天
- SLO：99.9%
- 错误预算：0.0999 (100–99.9)%
- 30 天总请求数：10000
- 允许的误差请求数：9.99 (10000 * 0.0999 / 100)

时间窗口

- 1w (7d)
- 4w(28d)
- 30d

告警级别


- page
- ticket

燃烧率

燃烧率	100%错误预算燃烧时间
1	30天
2	15天
60	12小时
1080	40分钟

SLO 告警指导思想-MWMMR

- **Precision**: 准确率
- **Recall**: 召回率 (故障漏过未告警)
- **Detection time**: 投递延迟
- **Reset time**: 告警重置时长



Severity	Long window	Short window	Burn rate	Error budget consumed
Page	1 hour	5 minutes	14.4	2%
Page	6 hours	30 minutes	6	5%
Ticket	3 days	6 hours	1	10%

```
expr: (  
    job:slo_errors_per_request:ratio_rate1h{job="myjob"} > (14.4*0.001)  
    and  
    job:slo_errors_per_request:ratio_rate5m{job="myjob"} > (14.4*0.001)  
)  
or  
(  
    job:slo_errors_per_request:ratio_rate6h{job="myjob"} > (6*0.001)  
    and  
    job:slo_errors_per_request:ratio_rate30m{job="myjob"} > (6*0.001)  
)  
severity: page
```


基于 Prometheus SLO 告警基础和挑战

开箱即用的 record and alert rule

```
# record
- record: code:prometheus_http_requests_total:sum
  expr: sum by (code) (prometheus_http_requests_total)

# alert
- alert: HighRequestLatency
  expr: job:request_latency_seconds:mean5m{job="myjob"} > 0.5
  labels:
    severity: page
```

加载和热更新

- 默认从本地文件加载。
- `kill -hub pid`
- `curl -X POST/PUT http://localhost:9090/-/reload`

Prometheus 告警基础

与时间窗口相关的多个 SLI rules

- `slo_errors_per_request:ratio_rate5m`
- `slo_errors_per_request:ratio_rate30m`
-

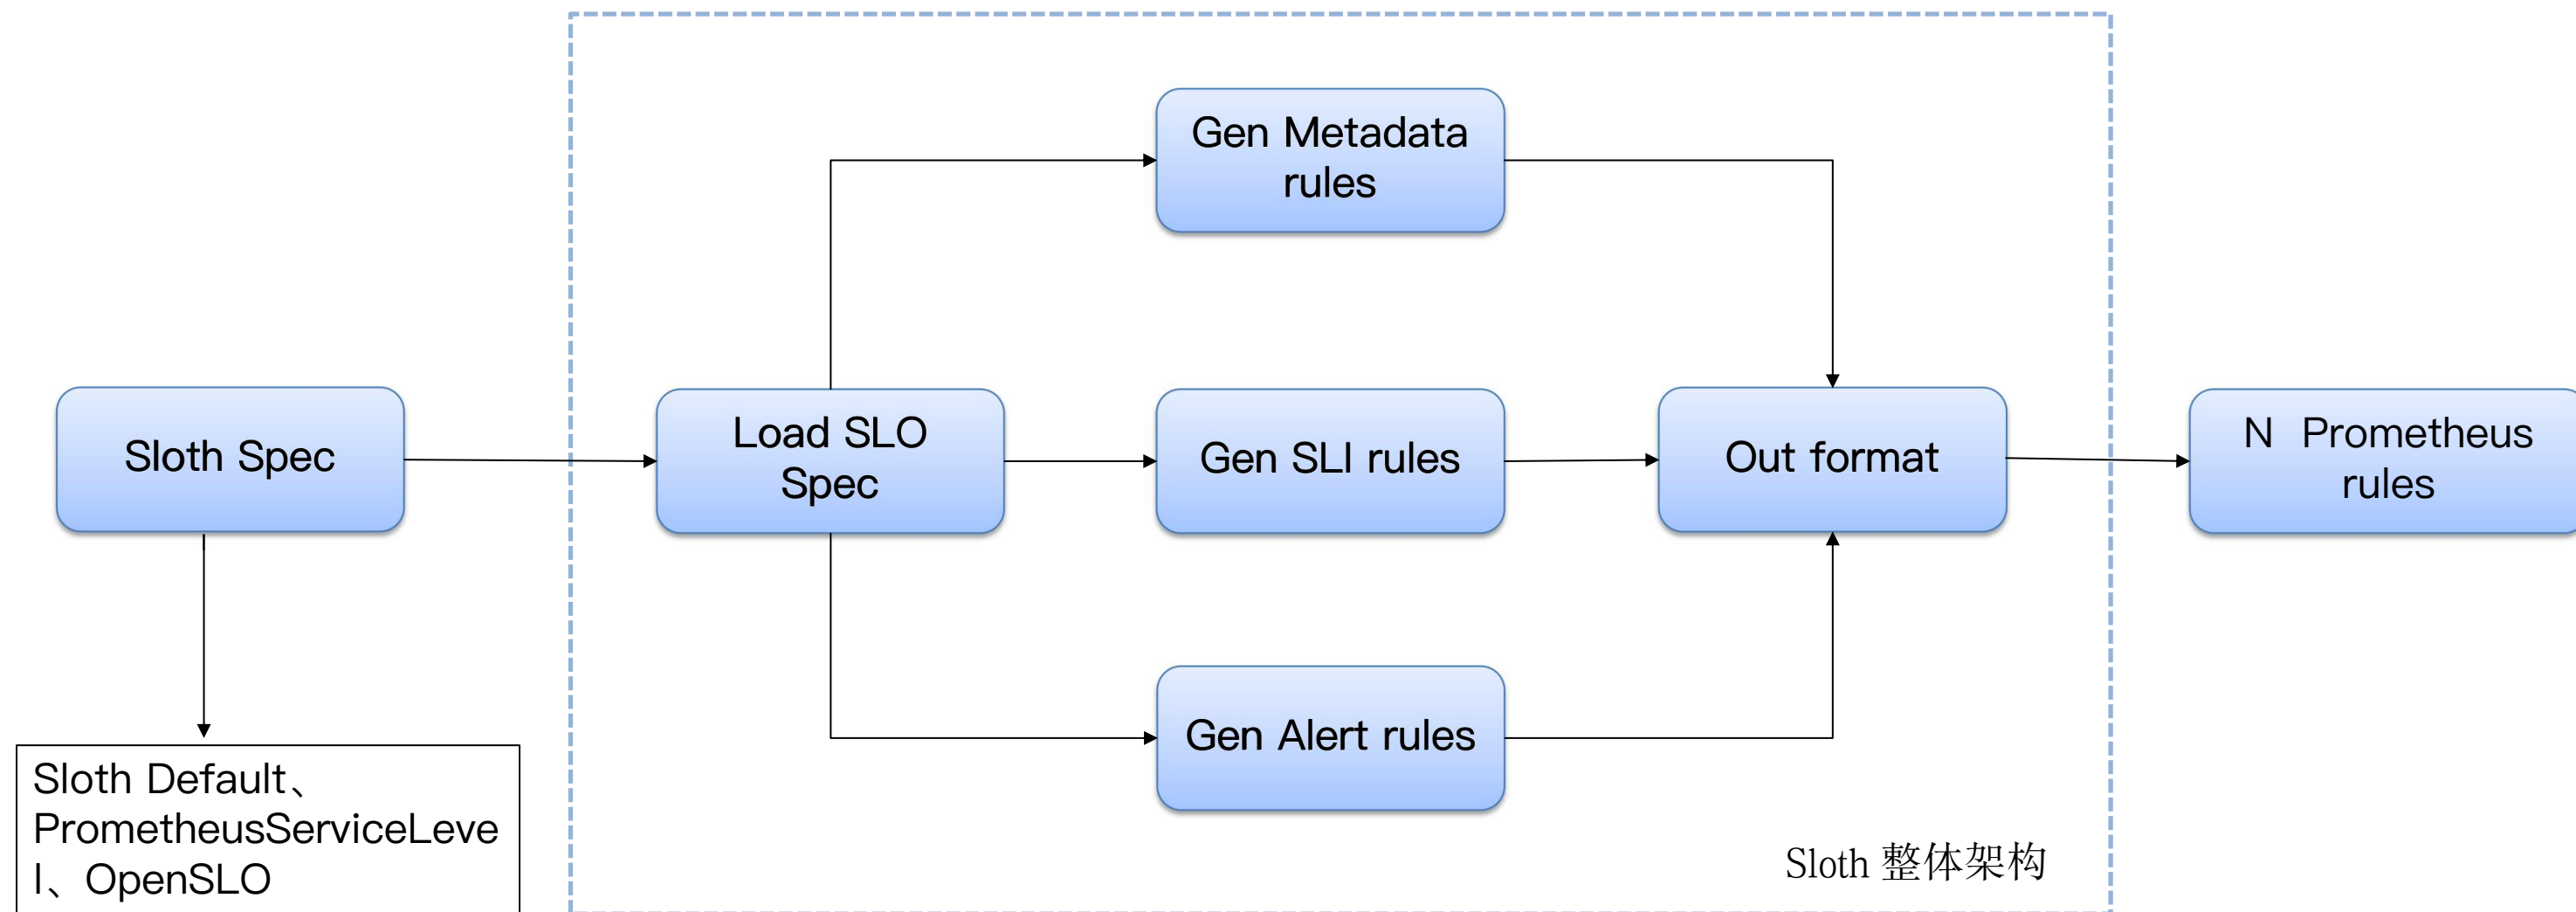
Alert rule 复杂, 需要考虑不同时间窗口和告警级别

```
expr: (
  job:slo_errors_per_request:ratio_rate1h{job="myjob"} > (14.4*0.001)
  and
  job:slo_errors_per_request:ratio_rate5m{job="myjob"} > (14.4*0.001)
)
or
(
  job:slo_errors_per_request:ratio_rate6h{job="myjob"} > (6*0.001)
  and
  job:slo_errors_per_request:ratio_rate30m{job="myjob"} > (6*0.001)
)
severity: page
```

Prometheus SLO 告警挑战

开源项目 sloth 简介

[Sloth](#) 是一个简单易用的 Prometheus SLO 自动生成器、支持 命令行和 K8s Controller 两种使用方式，支持自定义告警窗口配置、提供开箱即用的 Grafana 看板。



Sloth SLO 配置

Default sloth spec

```
version: "prometheus/v1"
service: "myservice" ← 以 service 来组织, 包含多个 slos
labels:
  owner: "myteam"
  repo: "myorg/myervice" } 公用标签
  tier: "2"
slos:
  # We allow failing (5xx and 429) 1 request every 100
  - name: "requests-availability"
    objective: 99.9
    description: "Common SLO based on availability"
    sli:
      events:
        error_query: sum(rate(http_request_duration_seconds_5xx{service="myservice"}))
        total_query: sum(rate(http_request_duration_seconds{service="myservice"}))
    alerting:
      name: MyServiceHighErrorRate
      labels:
        category: "availability" } Alert 通用标签
      annotations:
        # Overwrite default Sloth SLO alert summary
        summary: "High error rate on 'myservice' requests"
      page_alert:
        labels:
          severity: pageteam
          routing_key: myteam } Page 级别告警标签
      ticket_alert:
        labels:
          severity: "slack"
          slack_channel: "#alerts-myteam" } Ticket 级别告警标签
```

K8s CRD

```
apiVersion: sloth.slok.dev/v1
kind: PrometheusServiceLevel
metadata:
  name: sloth-slo-my-service
  namespace: monitoring
spec:
  service: "myservice"
  labels:
    owner: "myteam"
    repo: "myorg/myervice"
    tier: "2"
  slos:
    - name: "requests-availability"
      objective: 99.9
      description: "Common SLO based on availability"
      sli:
        events:
          errorQuery: sum(rate(http_request_duration_seconds_5xx{service="myservice"}))
          totalQuery: sum(rate(http_request_duration_seconds{service="myservice"}))
      alerting:
        name: MyServiceHighErrorRate
        labels:
          category: "availability"
        annotations:
          summary: "High error rate on 'myservice' requests"
        pageAlert:
          labels:
            severity: pageteam
            routing_key: myteam
        ticketAlert:
          labels:
            severity: "slack"
            slack_channel: "#alerts-myteam"
```

Open SLO

```
apiVersion: openslo/v1alpha
kind: SLO
metadata:
  name: sloth-slo-my-service
  displayName: Requests Availability
spec:
  service: my-service
  description: "Common SLO based on availability"
  budgetingMethod: Occurrences
  objectives:
    - ratioMetrics:
        good:
          source: prometheus
          queryType: promql
          query: sum(rate(http_request_duration_seconds{service="my-service"}))
        total:
          source: prometheus
          queryType: promql
          query: sum(rate(http_request_duration_seconds{service="my-service"}))
        target: 0.999
  timeWindows:
    - count: 30
      unit: Day
```

Sloth AlertWindows 配置

```
apiVersion: "sloth.slok.dev/v1"
kind: "AlertWindows"
spec:
```

```
  sloPeriod: 30d  —————→ SLO 时间周期
  page:         —————→ 不同告警级别
```

```
    quick:
```

```
      errorBudgetPercent: 2
      shortWindow: 5m
      longWindow: 1h
```

```
    slow:
```

```
      errorBudgetPercent: 5
      shortWindow: 30m
      longWindow: 6h
```

不同错误预算/燃烧率

多窗口

```
  ticket:
```

```
    quick:
```

```
      errorBudgetPercent: 10
      shortWindow: 2h
      longWindow: 1d
```

```
    slow:
```

```
      errorBudgetPercent: 10
      shortWindow: 6h
      longWindow: 3d
```

```
apiVersion: sloth.slok.dev/v1
kind: AlertWindows
spec:
```

```
  sloPeriod: 7d
```

```
  page:
```

```
    quick:
```

```
      errorBudgetPercent: 8
      shortWindow: 5m
      longWindow: 1h
```

```
    slow:
```

```
      errorBudgetPercent: 12.5
      shortWindow: 30m
      longWindow: 6h
```

```
  ticket:
```

```
    quick:
```

```
      errorBudgetPercent: 20
      shortWindow: 2h
      longWindow: 1d
```

```
    slow:
```

```
      errorBudgetPercent: 42
      shortWindow: 6h
      longWindow: 3d
```

AlertWindows 主要用于MWMR 告警配置，sloth 默认包含了 google-30d 和 google-28d 两个配置，可自定义。

<https://github.com/slok/sloth/tree/main/internal/alert/window>

Sloth CLI

```
sloth generate -i slos -o rules --slo-period-windows-path=./windows --default-slo-period="30d"
```

```
INFO[0000] SLI plugins loaded
INFO[0000] Using custom slo period windows catalog
INFO[0000] SLO period windows loaded
INFO[0000] Generating from Prometheus spec
INFO[0000] Multiwindow-multiburn alerts generated
s.Service version=v0.11.0 window=30d
INFO[0000] SLI recording rules generated
rometheus.Service version=v0.11.0 window=30d
INFO[0000] Metadata recording rules generated
rometheus.Service version=v0.11.0 window=30d
INFO[0000] SLO alert rules generated
rometheus.Service version=v0.11.0 window=30d
INFO[0000] Prometheus rules written
window=30d

plugins=0 svc=storage.FileSLIPlugin version=v0.11.0 window=30d
svc=alert.WindowsRepo version=v0.11.0 window=30d
svc=alert.WindowsRepo version=v0.11.0 window=30d windows=2
version=v0.11.0 window=30d
out=rules slo=myservice-requests-availability svc=generate.prometheu
out=rules rules=8 slo=myservice-requests-availability svc=generate.p
out=rules rules=7 slo=myservice-requests-availability svc=generate.p
out=rules rules=2 slo=myservice-requests-availability svc=generate.p
format=yaml groups=3 out=rules svc=storage.IOWriter version=v0.11.0
```

```
sloth generate -i slos -o rules --slo-period-windows-path=./windows --default-slo-period="28d"
```

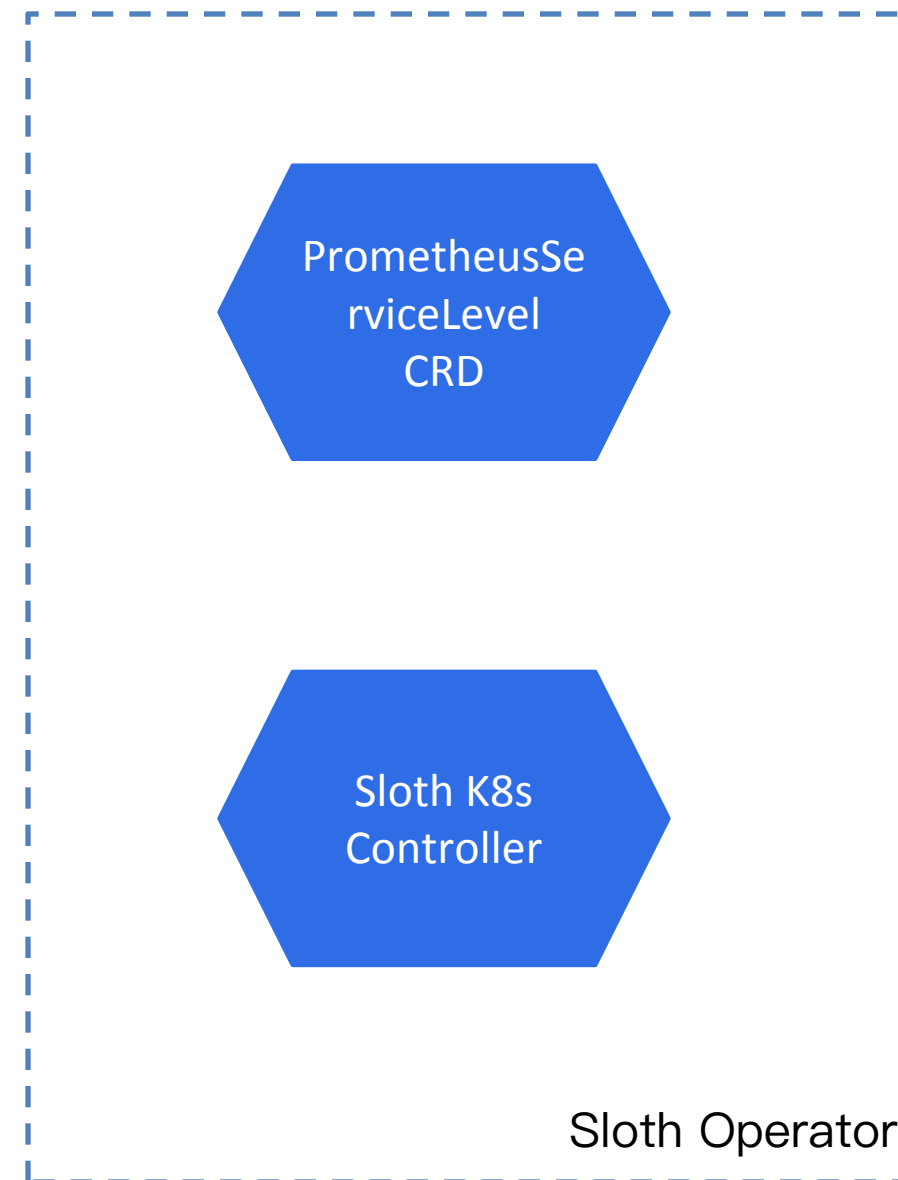
```
INFO[0000] SLI plugins loaded
INFO[0000] Using custom slo period windows catalog
INFO[0000] SLO period windows loaded
error: "generate" command failed: invalid default slo period: window period 672h0m0s missing
```

- 支持单个文件和目录批量生成。
- slo-period-windows 会覆盖默认 alert windows 配置
- default-slo-period 对应 alert windows 不存在，会报错

Sloth 与 K8s



Prometheus Operator



部署 sloth operator

```
kubectl apply -f  
https://raw.githubusercontent.com/slok/sloth/main/  
pkg/kubernetes/gen/crd/sloth.slok.dev\_prometheusservicelevels.ya  
ml
```

kubectl apply -f

```
https://raw.githubusercontent.com/slok/sloth/main/  
deploy/kubernetes/raw/sloth.yaml
```

部署 sloth SLO

kubectl apply -f

```
https://raw.githubusercontent.com/slok/sloth/main/  
examples/k8s-getting-started.yml
```

查看生成的 slos 和 prometheus rules

```
kubectl -n monitoring get slos
```

```
kubectl -n monitoring get prometheusrules
```


生成 Prometheus rules 详解

sli rules

- 8 个 record rules
- slo:sli_error:ratio_rate5m (30m、1h、2h、6h ...)

State	Name
Recording rule	slo:sli_error:ratio_rate5m
See graph	
Labels	owner=myteam repo=myorg/myservice sloth_id=myservice-requests-availability
Expression	<code>(sum(rate(http_request_duration_seconds_count{code=~\"[0-9]{1,3}\",sloth_id=\"myservice-requests-availability\",sloth_service=\"myservice\",sloth_slo=\"requests-availability\"}[5m])) / sum(rate(http_request_duration_seconds_count{job=\"myservice\",sloth_id=\"myservice-requests-availability\",sloth_service=\"myservice\",sloth_slo=\"requests-availability\"}[5m])))</code>
Recording rule	slo:sli_error:ratio_rate30m
Recording rule	slo:sli_error:ratio_rate1h
Recording rule	slo:sli_error:ratio_rate2h
Recording rule	slo:sli_error:ratio_rate6h
Recording rule	slo:sli_error:ratio_rate1d
Recording rule	slo:sli_error:ratio_rate3d
Recording rule	slo:sli_error:ratio_rate30d

metadata rules

- 6 个 record rules
- 包含了 SLO 目标、错误预算、时间周期等

State	Name
Recording rule	slo:objective:ratio
See graph	
Labels	owner=myteam repo=myorg/myservice sloth_id=myservice-requests-availability
Expression	<code>vector(0.9990000000000001)</code>
Recording rule	slo:error_budget:ratio
Recording rule	slo:time_period:days
Recording rule	slo:current_burn_rate:ratio
Recording rule	slo:period_burn_rate:ratio
Recording rule	slo:period_error_budget_remaining:ratio
Recording rule	sloth_slo_info

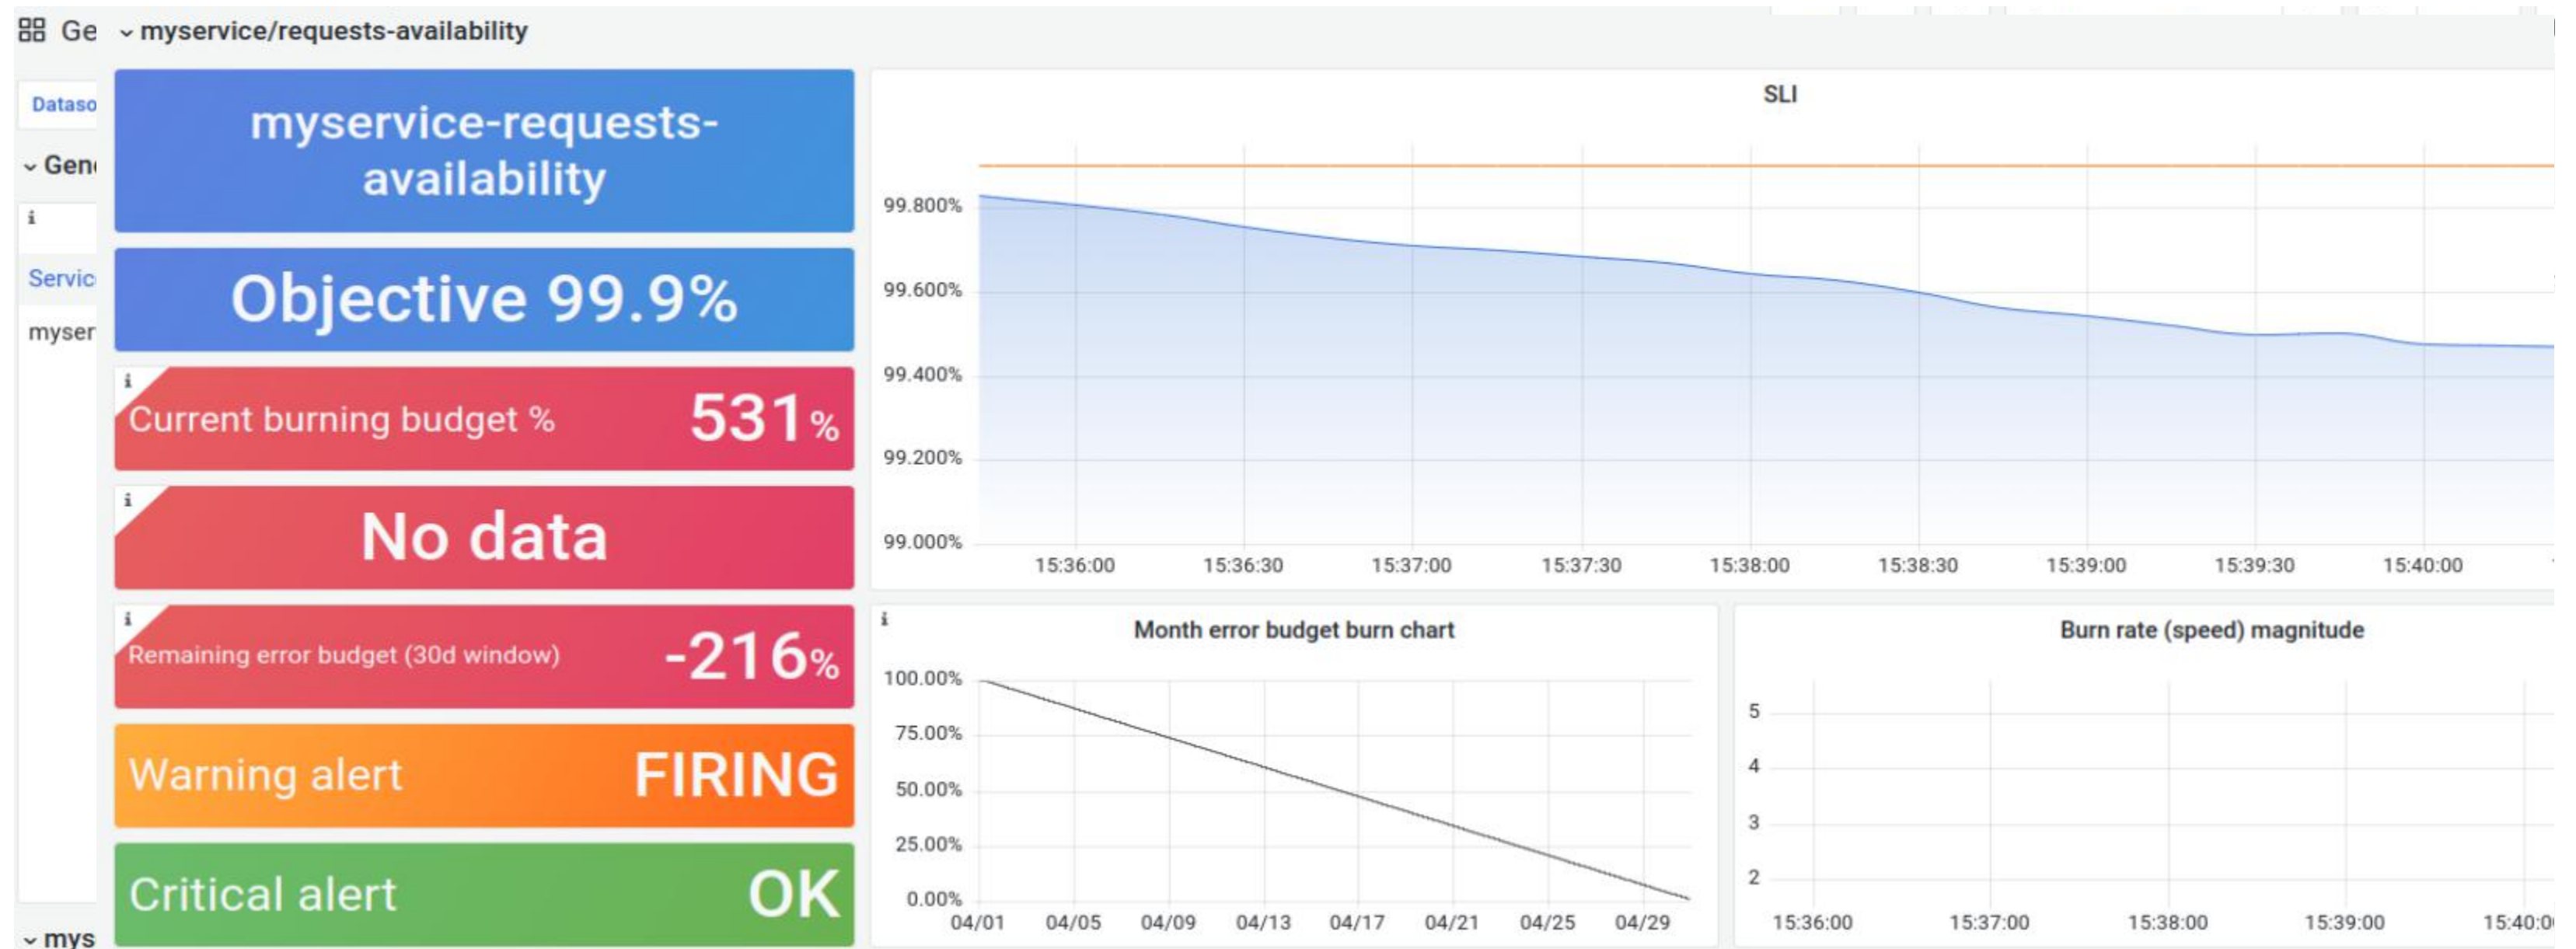
alert rules

- 2 个 alert rules
- 支持 MWMMR

State	Name	Health	Summary
Normal	MyServiceHighErrorRate	ok	High error rate on 'myservice' requests resp
Normal	MyServiceHighErrorRate	ok	High error rate on 'myservice' requests resp
See graph			
Labels	category=availability severity=slack slack_channel=#alerts-myservice sloth_severity=ticket		
Expression	<code>(max without (sloth_window) (slo:sli_error:ratio_rate2h{sloth_id="myservice-requests-availability",sloth_service="myservice",sloth_slo="requests-availability"} > {1 * 0.8000000000000001}) and max without (sloth_window) (slo:sli_error:ratio_rate1d{sloth_id="myservice-requests-availability",sloth_service="myservice",sloth_slo="requests-availability"} > {1 * 0.8000000000000001}) or (max without (sloth_window) (slo:sli_error:ratio_rate6h{sloth_id="myservice-requests-availability",sloth_service="myservice",sloth_slo="requests-availability"} > {1 * 0.8000000000000001}) and max without (sloth_window) (slo:sli_error:ratio_rate3d{sloth_id="myservice-requests-availability",sloth_service="myservice",sloth_slo="requests-availability"} > {1 * 0.8000000000000001})))</code>		
Summary	High error rate on 'myservice' requests responses		
Title	{label:sloth_service} {label:sloth_slo} SLO error budget burn rate is too fast.		
Matching	Search by label State		


Sloth Dashboard-Details

Template ID : 14348



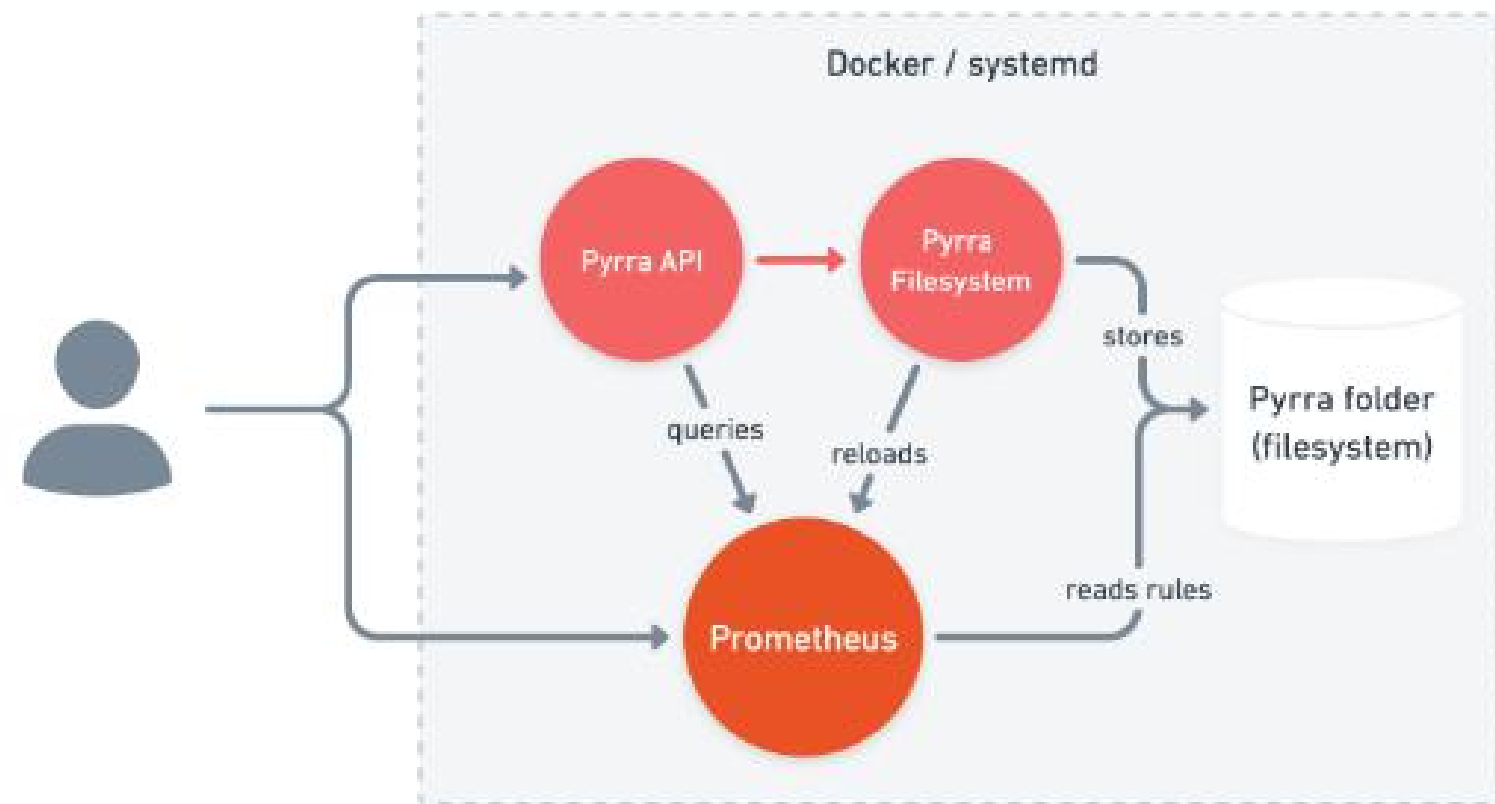
Sloth Dashboard-High Level

Template ID : 14348

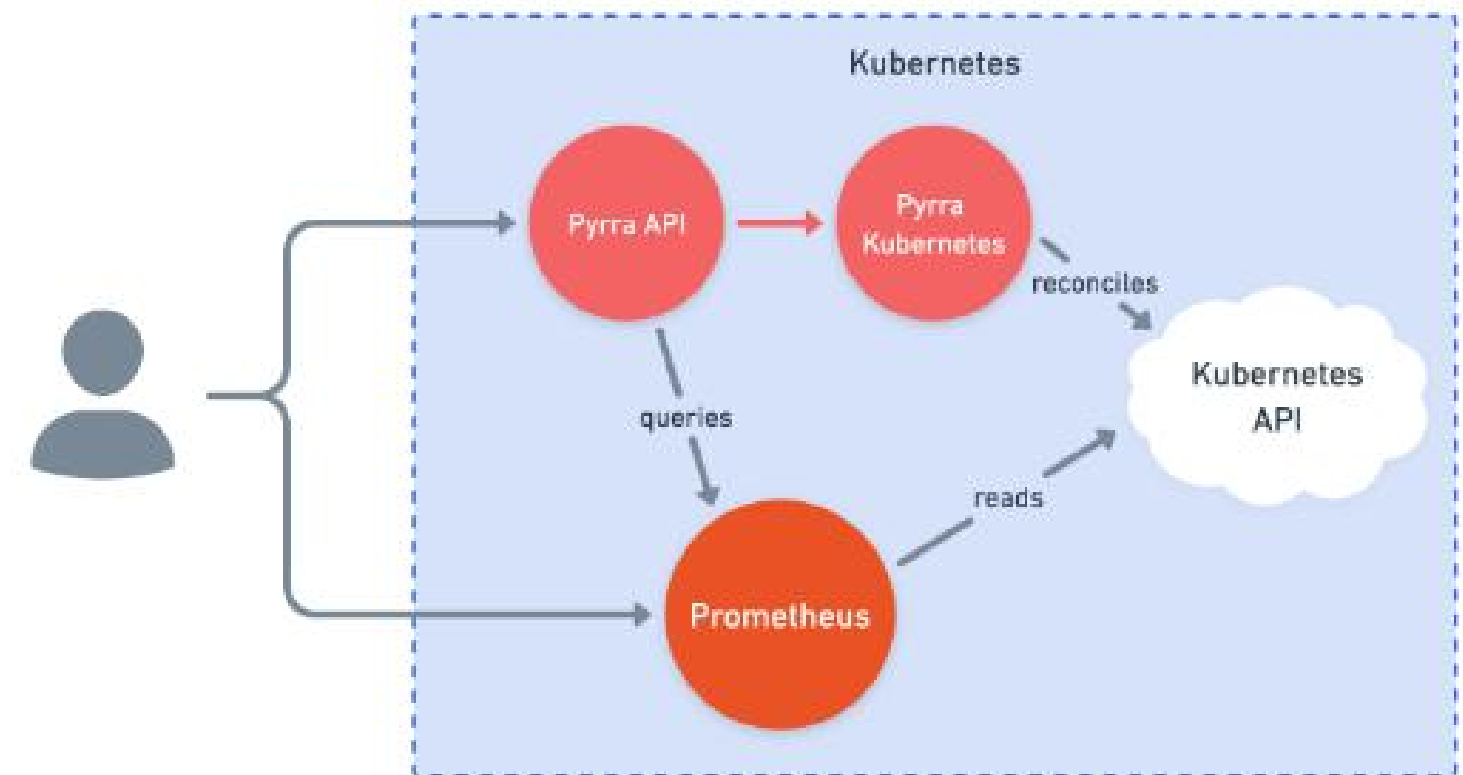


开源项目 Pyrra 介绍

[Pyrra](#) 是另外一个 Prometheus SLO 生成器、支持 filesystem 和 kubernetes 两种模式，提供 UI 界面和 Grafana 模板作为可视化。



Pyrra with filesystem



Pyrra with kubernetes

Pyrra SLO 配置

```
apiVersion: pyrra.dev/v1alpha1
kind: ServiceLevelObjective
metadata:
  name: pyrra-api-errors
  namespace: monitoring
  labels:
    prometheus: k8s
    role: alert-rules
    pyrra.dev/team: operations # Any labels prefixed with 'pyrra.dev/' will be propaga
spec:
  target: "99"
  window: 2w
  description: Pyrra's API requests and response errors over time grouped by route.
  indicator:
    ratio:
      errors:
        metric: http_requests_total{job="pyrra",code=~"5.."}
      total:
        metric: http_requests_total{job="pyrra"}
    grouping:
      - route
```

sli 和 alert rules 标签

SLO 目标值

时间窗口

Sli 分组

- 统一配置格式，无论是 K8s 还是 filesystem
- 一个 spec 定义一个 slo
- Slo 通过标签进行过滤，缺少 service 这层概念，按照配置文件名进行组织
- 不支持自定义告警窗口配置

Filesystem 和 K8s 模式使用

主要使用 pyrra filesystem 命令

```
/ # pyrra filesystem --help
Usage: pyrra filesystem

Runs Pyrra's filesystem operator and backend for the API.

Flags:
  -h, --help                Show context-sensitive help.
  --config-files="/etc/pyrra/*.yaml"
                             The folder where Pyrra's configuration files are located.
  --prometheus-url=http://localhost:9090
                             The URL to the Prometheus server.
  --prometheus-folder="/etc/prometheus/pyrra/"
                             The folder where Prometheus rules are stored.
  --generic-rules            Enabled generic rules.
```

- 与 Prometheus server 处于同一实例，pyrra 加载目录配置，生成 rules 能够确保被 Prometheus server 加载。
- 通过配置的 prometheus-url ,进行 hot reload让 Prometheus 重新加载配置。

Filesystem 模式

部署 pyrra operator

```
kubectl apply -f ./config/crd/bases/pyrra.dev_servicelevelobjectives.yaml
kubectl apply -f ./config/rbac/role.yaml
kubectl apply -f ./config/api.yaml
kubectl apply -f ./config/kubernetes.yaml
```

部署 pyrra slos

```
kubectl apply -f ./examples/kubernetes/slos/
```

查看生成的 slos 和 prometheus rules

```
kubectl -n monitoring get servicelevelobjectives
kubectl -n monitoring get prometheusrules
```

K8s 模式

Pyrra 生成 Prometheus rules 详解

单个 SLO 包括三个 Prometheus 告警分组，分别为 xxx-availability、xxx-availability-generic、xxx-availability-increase.

xxx-availability: 包含 7 个 SLI record rules 和 4个 MWMR alert rules。

>	Recording rule	http_request_duration_seconds:burnrate5m
>	Recording rule	http_request_duration_seconds:burnrate32m
>	Recording rule	http_request_duration_seconds:burnrate1h4m
>	Recording rule	http_request_duration_seconds:burnrate2h9m
>	Recording rule	http_request_duration_seconds:burnrate6h26m
>	Recording rule	http_request_duration_seconds:burnrate1d1h43m
>	Recording rule	http_request_duration_seconds:burnrate4d6h51m
>	Normal	ErrorBudgetBurn
>	Normal	ErrorBudgetBurn
>	Normal	ErrorBudgetBurn
>	Normal	ErrorBudgetBurn

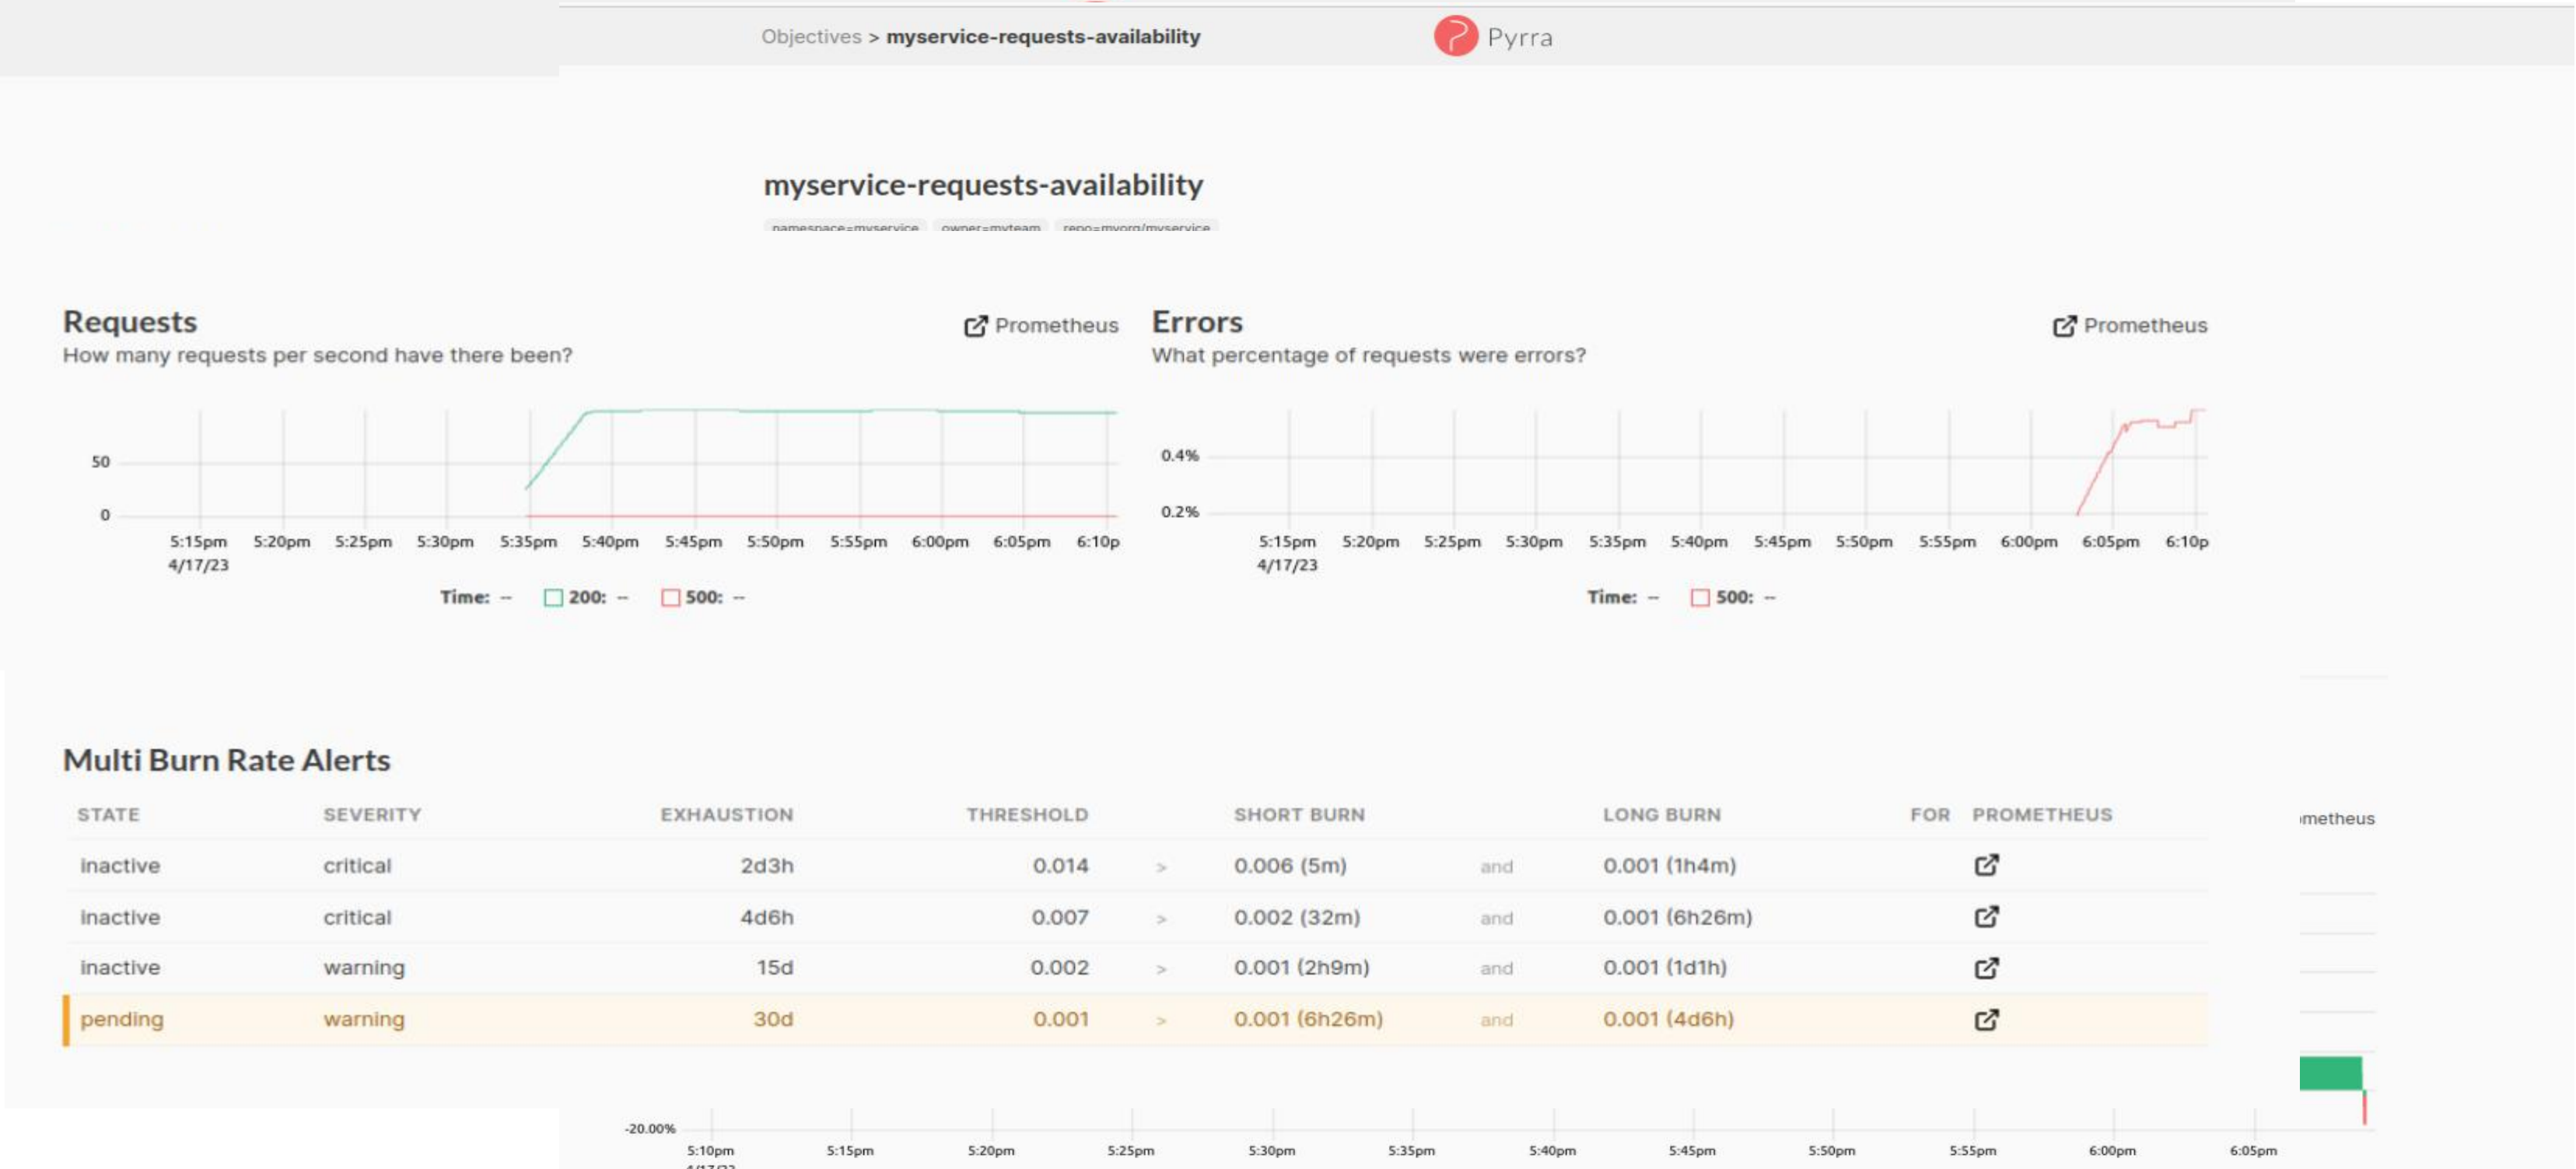
xxx-availability-generic： 5 个 record rules， 主要用于看板的 high level 汇总统计（ RED ）。

State	Name
➤ Recording rule	pyrra_objective
➤ Recording rule	pyrra_window
➤ Recording rule	pyrra_availability
➤ Recording rule	pyrra_requests_total
➤ Recording rule	pyrra_errors_total

xxx-availability-increase： 1 个 record rule 和 1个 alert rule， 主要用于统计整个窗口周期总请求数和 no data 告警。

>	Recording rule	http_request_duration_seconds:increase30d
>	Normal	SLOMetricAbsent

Pyrra Dashboard-API



Multi Burn Rate Alerts

STATE	SEVERITY	EXHAUSTION	THRESHOLD		SHORT BURN		LONG BURN	FOR	PROMETHEUS
inactive	critical	2d3h	0.014	>	0.006 (5m)	and	0.001 (1h4m)		
inactive	critical	4d6h	0.007	>	0.002 (32m)	and	0.001 (6h26m)		
inactive	warning	15d	0.002	>	0.001 (2h9m)	and	0.001 (1d1h)		
pending	warning	30d	0.001	>	0.001 (6h26m)	and	0.001 (4d6h)		

5:10pm

5:15pm

5:20pm

5:25pm

5:30pm

5:35pm

5:40pm

5:45pm

5:50pm

5:55pm

6:00pm

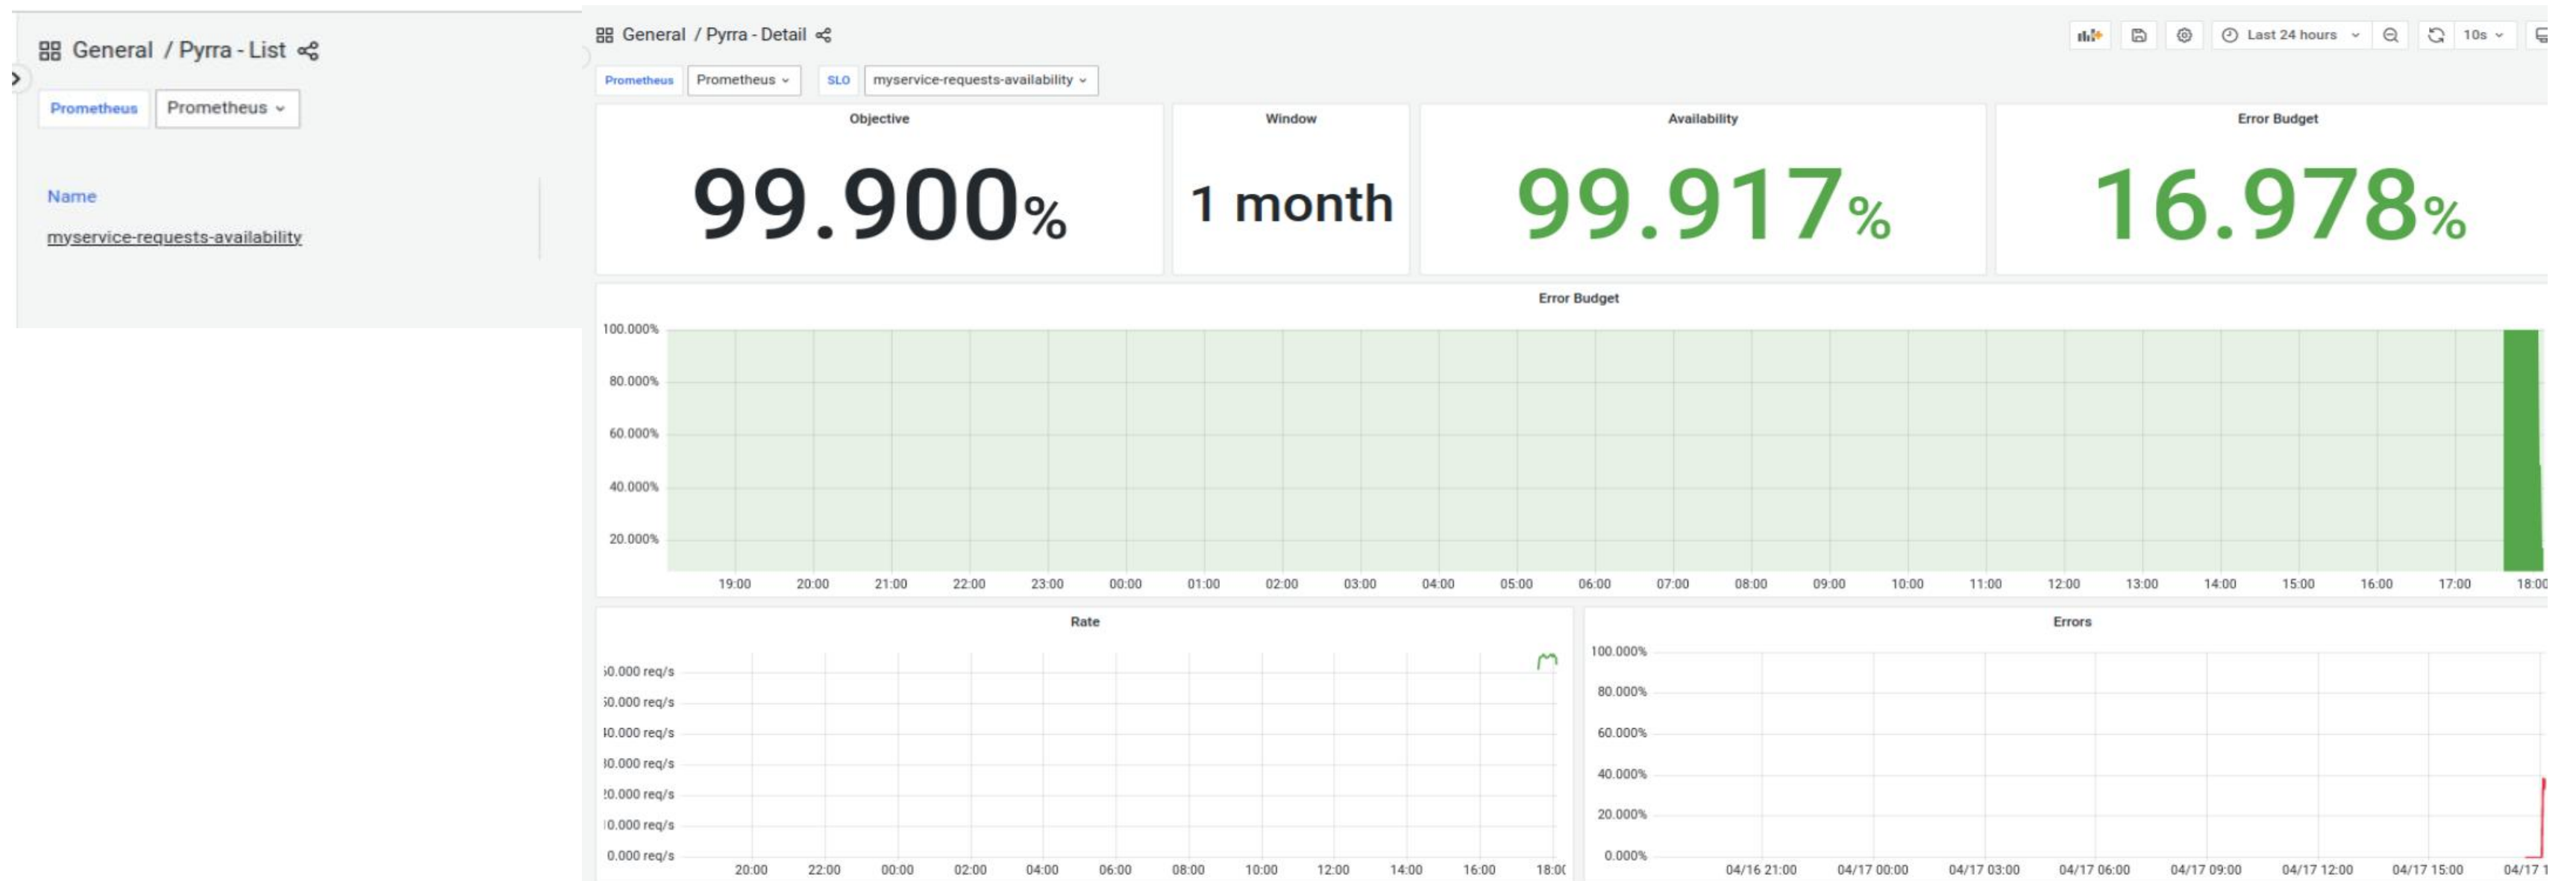
6:05pm

4/17/23

-20.00%

Pyrra Dashboard-Grafana

Grafana 模板 json <https://github.com/pyrra-dev/pyrra/tree/main/examples/grafana>



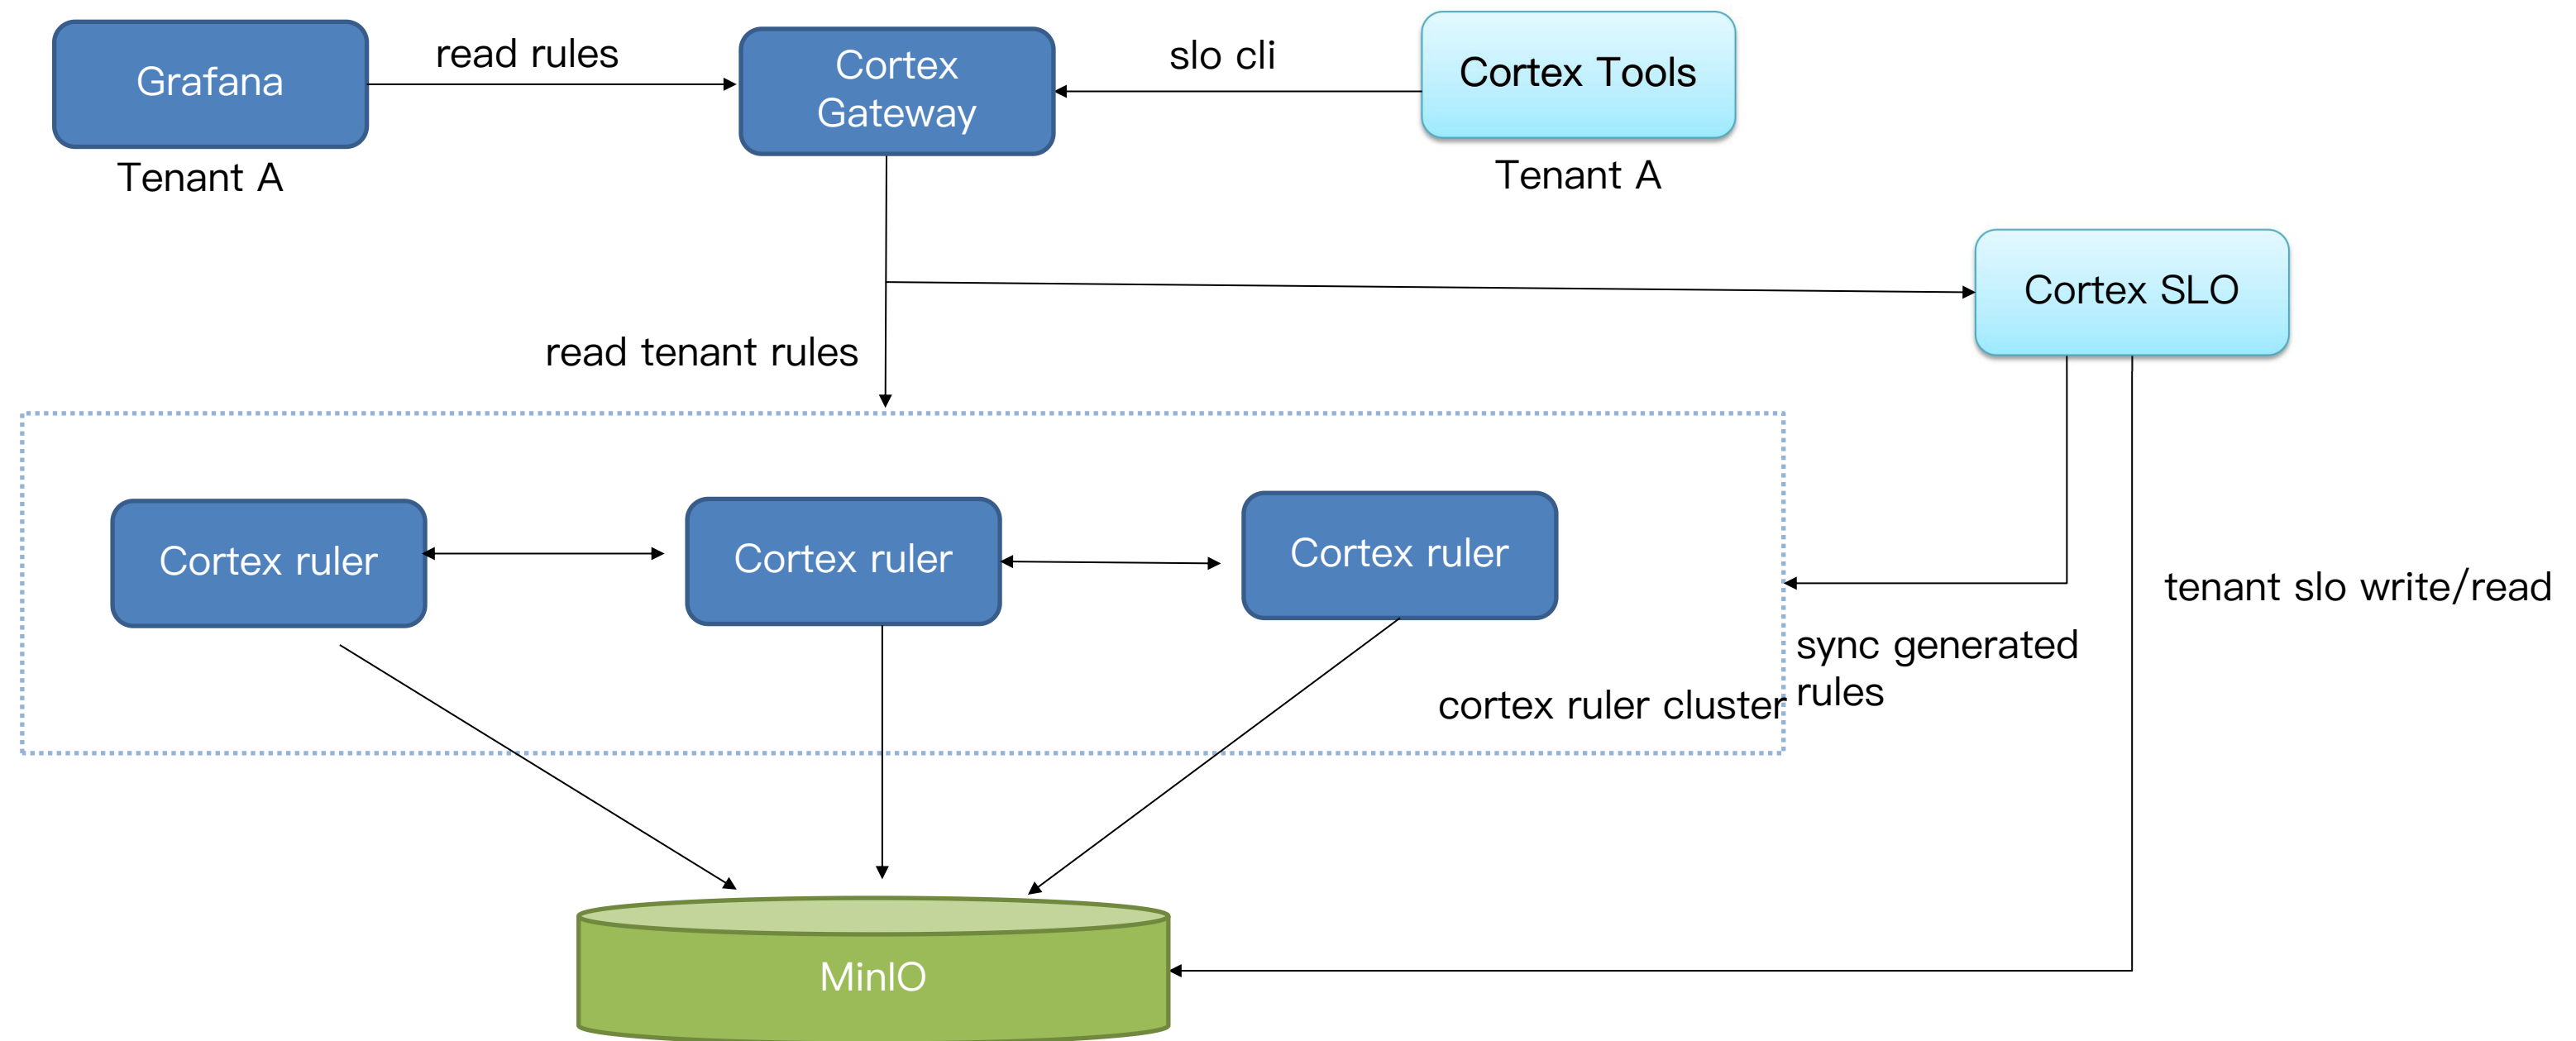
Sloth Vs Pyrra

	Sloth	Pyrra
Github Star	1.5k	800+
K8s	支持	支持
Filesystem	不支持	支持
CLI	支持	不支持
OpenSLO	支持	不支持
SLI 可读性	高	一般
Dashboard	Grafana	Grafana、Pyrra API

简单总结：

- 两个都是优秀的 Prometheus SLO generator，sloth 开源时间较早，协议支持广泛，pyrra 属于后起之秀，有自己的 dashboard。
- 因为 sloth 生成的 rules 可读性更强，如果有二开需求并直接使用 Grafana 作为看板，建议采用 sloth。

基于 Cortex 的多租户 SLO 服务构建



Cortex-tools 扩展

统一封装到 slos 子命令

```
1 $ ./cortextool slos --help
usage: cortextool slos [<flags>] <command> [<args> ...]

View & edit slos stored in cortex.

Flags:
  --help           Show context-sensitive help (also try --help
  --authToken=""   Authentication token for bearer token or JWT

Subcommands:
  slos list [<flags>]
    List the slos currently in the telnant.

  slos load [<flags>] <files>...
    Load a set of slos to a designated cortex endpoint.

  slos get [<flags>] <service>
    Retrieve the sevice slos from the cortex.

  slos delete [<flags>] <service>
    Remove the sevice slos from the cortex.

  slos list-windows [<flags>]
    List slo windows currently in the telnant.

  slos load-windows [<flags>] <files>...
    Load a set of windows to a designated cortex endpoint.

  slos get-windows [<flags>] [<windows>]
    Get windows to from remote store.

  slos delete-windows [<flags>] [<windows>]
    Get windows to from remote store.
```

配置 cortex-slo 地址以及租户信息

```
export CORTEX_ADDRESS=http://localhost:6666
export CORTEX_TENANT_ID=demo
```

#导入 和查询 windows

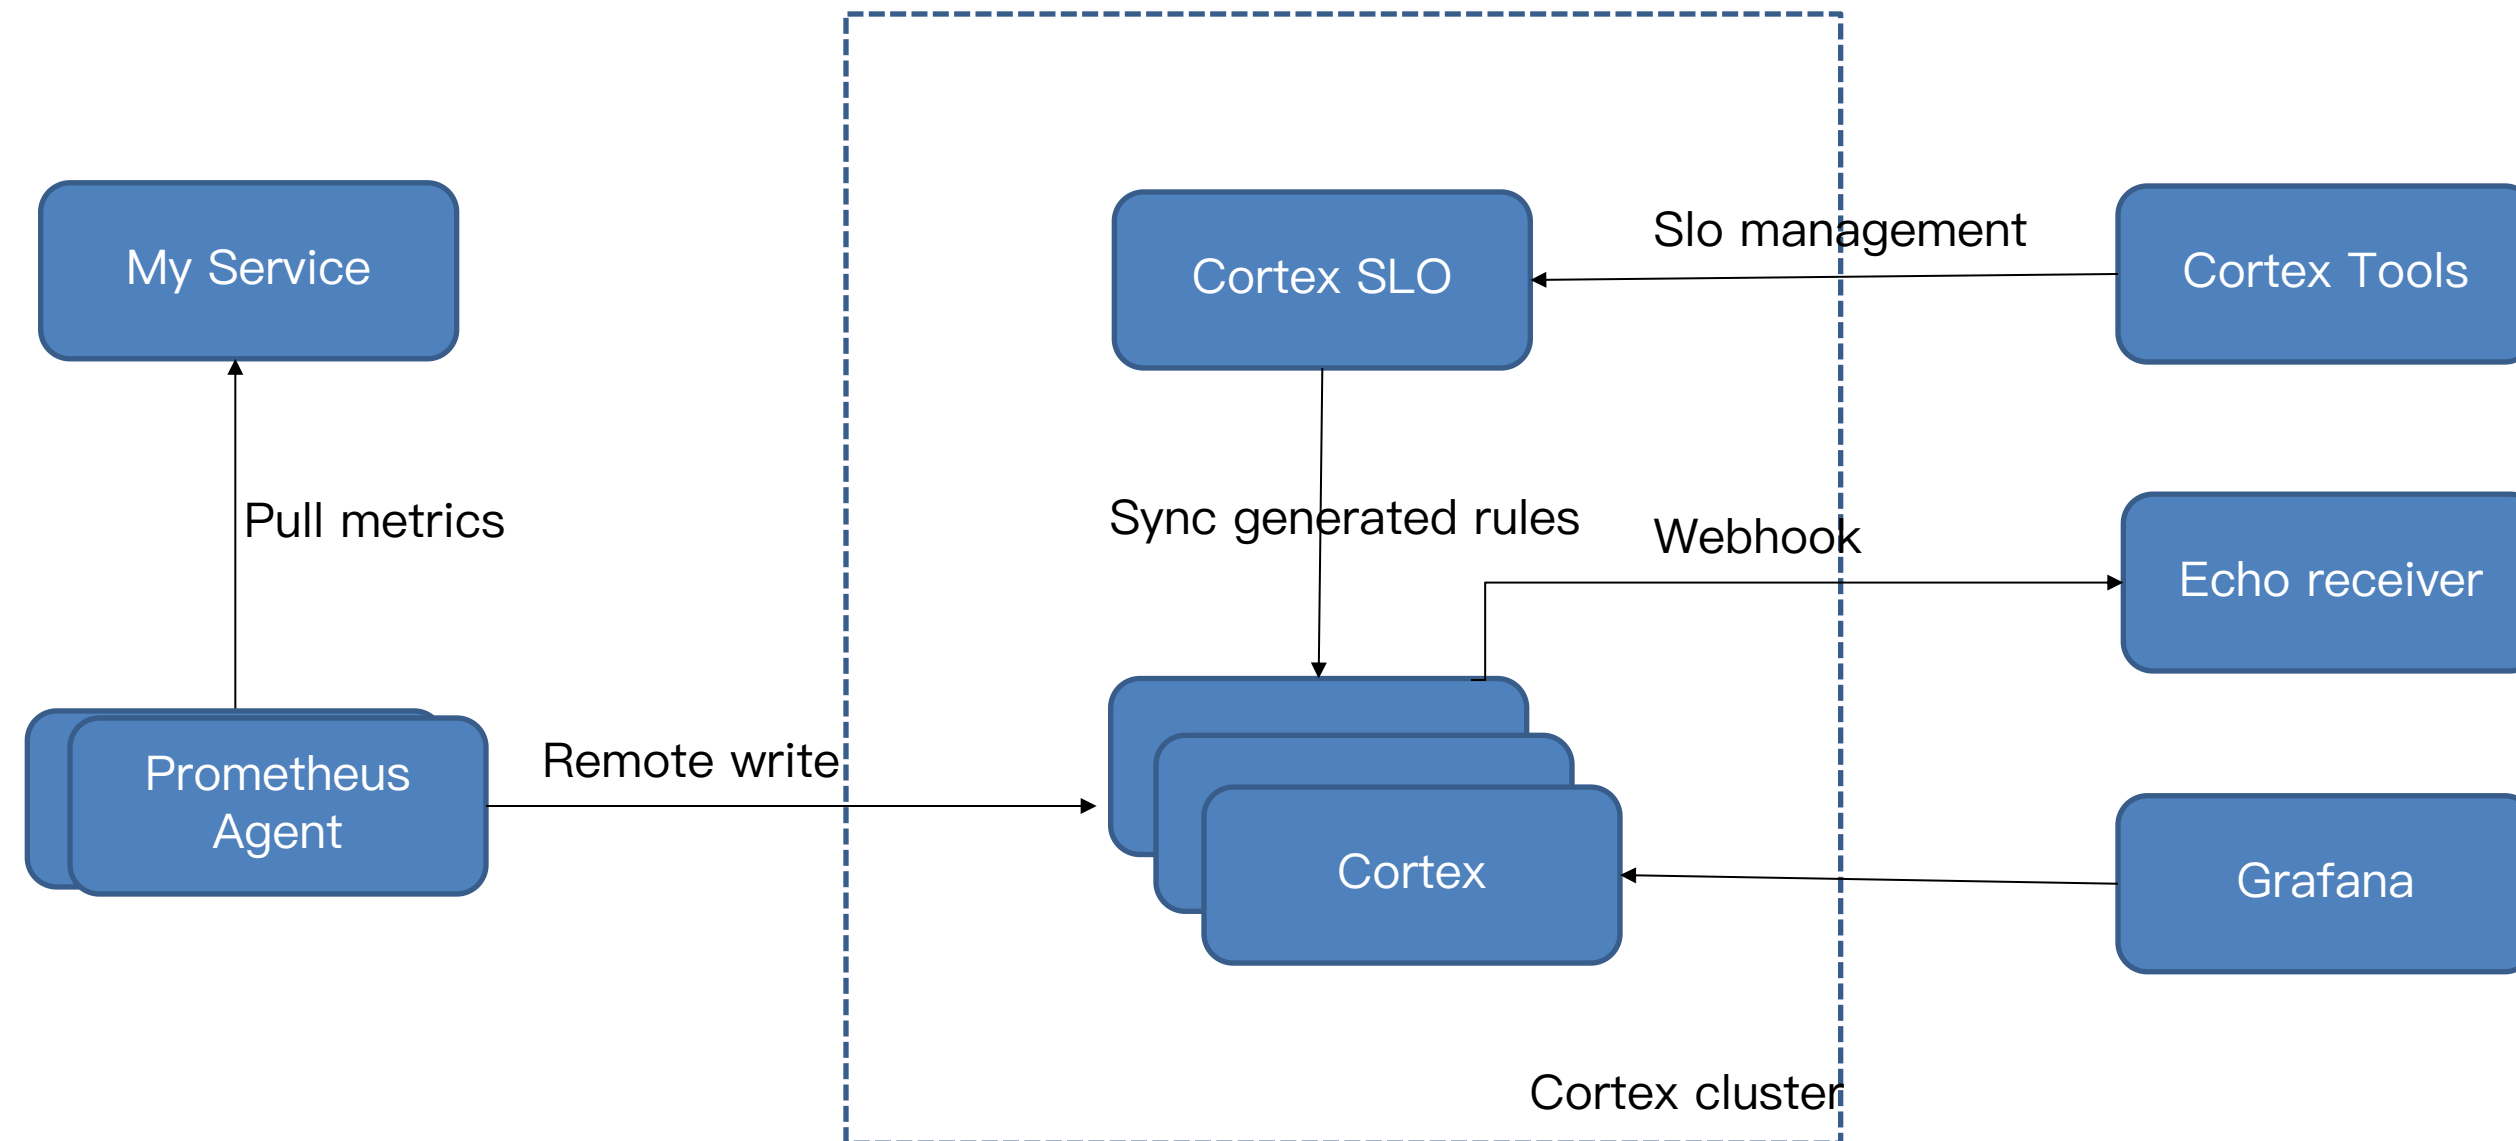
```
11:44 $ cortextool slos load-windows ./config/slos/windows/*
INFO[0000] 7d.yaml windows loaded
INFO[0000] google-30d.yaml windows loaded
✓ /home/service/workspace/tower/play-with-cortex-slo [main|+2...5]
1 $ cortextool slos list-windows

INFO[0000] Windows:
7d
google-30d
✓ /home/service/workspace/tower/play-with-cortex-slo [main|+2...5]
1 $ cortextool slos load ./config/slos/*.yaml --windows google-30d
INFO[0000] myservice.yaml slos loaded with google-30d alert windows
✓ /home/service/workspace/tower/play-with-cortex-slo [main|+2...5]
1 $ cortextool slos list

INFO[0000] Slos:
myservice
✓ /home/service/workspace/tower/play-with-cortex-slo [main|+2...5]
1 $ cortextool slos get myservice

INFO[0000] myservice Slos:
version: "prometheus/v1"
service: "myservice"
labels:
  owner: "myteam"
  repo: "myorg/myservice"
slos:
```


Demo with docker-compose



参考 GitHub 仓库 <https://github.com/grafanafans/play-with-cortex-slo>

资料

- Google SRE workbook <https://sre.google/workbook/alerting-on-slos>
- Sloth GitHub <https://github.com/slok/sloth>
- Pyrra GitHub <https://github.com/pyrra-dev/pyrra>
- Play with sloth <https://github.com/grafanafans/play-with-sloth>
- Play with pyrra <https://github.com/grafanafans/play-with-pyrra>



OBSERVABILITY
SUMMIT 2023
可观测性峰会 第1届

Thank you



关注我们获取更多云原生资讯